



Working Paper Series

Working Paper No. 26-02 • August 2026

Aid Cuts and the Politics of Refugee Hosting

Mae MacDonald

Aid Cuts and the Subnational Foundations of Refugee Cooperation

Mae MacDonald*

Sept 13, 2026

Word count: 7,925

Abstract

What happens to international cooperation when the transfers that sustain its domestic political foundations are withdrawn? I argue that cooperation with the international refugee regime rests partly on a subnational bargain: host communities bear concentrated costs of asylum, but foreign aid can offset these costs by providing jobs, services, and other benefits. Donor withdrawal should therefore weaken support where the refugee aid economy is concentrated, even if national public opinion remains stable. I examine this argument during the 2025 aid cuts in Kenya using two nationwide surveys, a high-frequency panel and vignette experiment near Kakuma refugee camp, and 52 interviews with refugee leaders. Support for refugee hosting declined in refugee-hosting regions relative to the rest of the country, and within the host community, the decline was concentrated among households most exposed to the loss of aid-related jobs or services. Yet national support remained stable. The findings show that international cooperation can appear stable nationally even as support erodes among the subnational constituencies on which its cooperation depends.

*Department of Political Science, Stanford University. Email: maemac@stanford.edu. ORCID: 0000-0002-9055-4810.

I am grateful to Jeremy Weinstein, Jens Hainmueller, James D. Fearon, Michael Allen, and Judith Goldstein for their invaluable advising and feedback. I thank Claire Adida, Sigrid Weber, Shelby Carvalho, Allison Grossman, Paige Hill, and participants in the 4th UNHCR–World Bank Joint Data Center Research Conference on Forced Displacement, 2025 American Political Science Association Annual Conference, U-Penn PDRI Forced Displacement Working Group, Stanford Immigration Policy Lab workshop, and Stanford IPE Workshop for helpful comments. I thank Mary Achieng, Michael Mwarange, Maggie Ireri, and the team at TIFA Research for outstanding work implementing the surveys, as well as Tolossa Asrat for excellent research assistance. I sincerely thank the participants in this research for sharing their experiences, particularly the refugee leaders who spoke with me during interviews. This project was supported by funding from the J-PAL/IPA Displaced Livelihoods Initiative and the Stanford King Center on Global Development.

1 Introduction

What happens to international cooperation when the transfers that sustain its domestic political foundations are withdrawn? A central insight in international political economy (IPE) is that interstate cooperation often depends on domestic political bargains. International agreements create winners and losers, and governments may need to compensate those who bear the costs to sustain support for cooperation. This logic is particularly well developed in research on international trade (e.g., Putnam, 1988; Ruggie, 1982; Broz, Frieden and Weymouth, 2021). Yet we know much less about the domestic political foundations of cooperation with the international refugee regime, which protects 42.5 million refugees, 71% of whom live in low- and middle-income countries (LMICs).

I posit that understanding cooperation with the refugee regime requires looking below the level of the state. Existing accounts characterize the regime as a “grand compromise,” an interstate bargain whereby host states grant refugees asylum in their territory, while donor states provide foreign aid to finance refugee protection (Betts, 2009; Cuéllar, 2006). This bargain helps solve a collective-action problem at the international level. But its implementation is highly unequal within host countries. Refugees are often concentrated in a small number of poor and politically marginalized border regions, particularly in LMICs where many live in camps rather than dispersing throughout the country.

This creates a second, domestic, political bargain. Refugee-hosting communities bear concentrated costs of asylum, including pressure on land and services, economic competition, and insecurity. But the aid economy that accompanies refugees can also generate concentrated local benefits through schools, health facilities, jobs, and broader economic development (Zhou, Grossman and Ge, 2023; Alix-Garcia et al., 2018; Kadigo and Maystadt, 2023; Zhou et al., 2025; Demir, Mayda and Maystadt, 2025; Baseler et al., 2025; MacDonald and Lichtenheld, 2025). When these benefits offset the costs of hosting, communities may support or even welcome refugee presence. When this subnational bargain breaks down, however, opposition in hosting communities can generate protests and conflict and increase pressure on governments to restrict refugee hosting.

A distinctive feature of this subnational bargain is that compensation comes not from the host government, as in the trade regime, but from donors and international organizations. Local support may therefore be vulnerable to donor withdrawal when assistance no longer offsets the concentrated costs of hosting. The abrupt contraction of foreign aid in 2025 provides a tragic test of this argument. Official development assistance from OECD countries fell by 23.1%, driven largely by the U.S. dismantling of USAID (OECD, 2026). Beyond the severe direct welfare consequences of these cuts (Ferreira da Silva et al., 2026; Kenny and Sandefur, 2025), the loss of aid-related ben-

efits may reduce host communities' support for refugees and increase pressure for their removal, weakening the subnational foundations of cooperation.

In this paper, I test the first link in this broader argument: whether donor withdrawal reduces support for refugee hosting among host communities, and whether this response is concentrated among citizens who personally lose aid-related benefits. I study this question in Kenya, a major refugee-hosting country and one of the largest recipients of U.S. assistance before the cuts, with features common to many hosting contexts: protracted displacement, large refugee camps, high aid dependence, and economically marginalized host regions. I draw on two nationwide phone surveys, a high-frequency face-to-face panel near one of Kenya's two major refugee camps that includes a survey experiment, and 52 panel interviews with refugee leaders, conducted before, during, and after the most severe phase of the 2025 aid cuts.

The analysis yields three findings. First, support for refugees in host communities depends in large part on the material benefits accompanying refugee aid, including access to refugee schools and health facilities, employment with refugees and NGOs, and business opportunities generated by refugee and aid-worker demand. Second, the aid cuts generated a localized backlash against refugee hosting. Citizens in refugee-hosting regions became less willing to support refugees' continued local presence relative to citizens elsewhere in Kenya. Within the host community, this decline was concentrated among households that directly lost aid-related jobs or services, rather than among nearby households that did not. Qualitative evidence links the decline to lost economic benefits and growing insecurity.

Third, this backlash did not extend nationwide. Support remained stable outside the refugee-hosting regions, producing convergence as initially higher support in hosting areas declined toward the national average. The political effects of donor withdrawal were therefore concentrated where its material effects were greatest, even as aggregate public opinion made support appear stable. Exploratory analyses also show high levels of empathy and humanitarian concern among host communities and the broader public, consistent with research suggesting that humanitarian concerns sustain support for refugees in LMICs (Newman et al., 2015; Pettigrew and Tropp, 2006; Ghosn, Braithwaite and Chu, 2019).

This study makes three contributions. First, it connects research in IPE on the domestic politics of international cooperation with work in comparative politics and development economics showing that refugee hosting and related aid can generate economic development for host countries. I argue that such material benefits can sustain local constituencies for international cooperation with the refugee regime. More broadly, the paper applies an IPE lens to the refugee regime, which has received comparatively little attention relative to regimes governing trade, finance, the environment, and human rights.

Second, the paper provides micro-level evidence that donor withdrawal can erode support for refugee hosting, with important policy implications. In Kenya, backlash was concentrated in refugee-hosting regions and among households that lost aid-related benefits. Historical experience suggests that such local opposition can matter nationally: in the 1990s in Kenya, for example, violence and opposition around coastal refugee camps contributed to the government closing the camps and requiring refugees to relocate or return to Somalia. Similar dynamics may be emerging elsewhere as cuts to refugee assistance affect LMICs across Africa, Asia, and the Middle East.¹

Third, the paper provides rare real-time evidence on the local consequences of a major international policy shock. I designed and launched the study soon after the cuts were announced and before their trajectory was clear, allowing me to observe attitudes during the sharpest contraction in assistance. High uncertainty, security risks, and restrictions on access made data collection especially challenging. I therefore combined quantitative and qualitative evidence across multiple methods and levels of analysis, triangulating the findings across sources. The study’s key empirical contribution is thus to capture attitudinal change as the cuts unfolded rather than reconstructing it retrospectively.

2 The Subnational Foundations of Cooperation with the Refugee Regime

The international refugee regime addresses a collective action problem: people fleeing persecution may seek asylum and safety elsewhere, but states have incentives to shift the costs of providing protection onto others (Suhrke, 1998; Betts, 2003). Built after the Second World War, scholars have characterized the regime as a “grand compromise” between the low- and middle-income countries that host most refugees and the high-income donor states that finance assistance through organizations such as UNHCR, WFP, and IOM (Cuéllar, 2006; Betts, 2009).² Scholarship on cooperation with the refugee regime has largely focused on this interstate bargain.

Yet this obscures how refugee protection is implemented within host states. National governments decide whether refugees may remain, but the consequences are often concentrated in poor and politically marginalized border communities, as refugees often settle or are accommodated in camps near the borders of the countries they fled (Anti

¹Major funding shortfalls have affected refugee assistance in Uganda (Ariba, 2026), Thailand (Regan et al., 2025), Bangladesh (Associated Press, 2026), Chad (Farge, 2026), Jordan (Jordan Times, 2026), and elsewhere. WFP projected a 34% funding decline, while UNHCR’s funding fell by 24% and the agency cut one-third of its workforce (UNHCR, 2026a; WFP, 2025a). Existing studies show that earlier reductions in WFP assistance harmed refugees, though they do not estimate effects on hosts (Grossman, Zhou and Carvalho, 2025; Sterck and Bruni, 2025).

²Although Europe and North America have increasingly received refugees (Arar, 2017), 71% of refugees remain in LMICs.

et al., 2025). Residents in these areas share land, markets, services, infrastructure, and security conditions with refugees. These communities can in turn shape national policy, as local opposition can generate conflict and pressure to remove refugees, while acceptance can sustain prolonged hosting. Highly exposed host communities may therefore be an important constituency for refugee cooperation alongside the governments and broader electorates emphasized in existing research.

While refugees can impose costs on these host communities (Verme and Schuettler, 2021; Ceritoglu et al., 2017; Salehyan and Gleditsch, 2006), refugee aid can offset them by bringing schools, health facilities, jobs, business demand, and infrastructure to otherwise marginalized regions (Alix-Garcia et al., 2018; International Finance Corporation, 2018; World Bank and UNHCR, 2021; Zhou, Grossman and Ge, 2023; Kadigo and Maystadt, 2023). These benefits may in turn increase host-community support for refugees (Baseler et al., 2025; Zhou et al., 2025; MacDonald and Lichtenheld, 2025; Demir, Mayda and Maystadt, 2025).

Unlike familiar accounts of domestic compensation in international political economy, however, much of this compensation is financed by external actors, particularly the U.S., which channels support through international organizations such as UNHCR, IOM, WFP, and UNRWA. This dependence on external financing makes the local foundations of refugee cooperation vulnerable to donor withdrawal. Aid cuts directly affect households employed by NGOs or refugees, residents who rely on camp-based health and education services, and businesses dependent on refugee and aid-worker demand. They may also increase the costs of hosting if declining assistance worsens refugee deprivation, competition over scarce resources, or insecurity.

I therefore hypothesize that the political effects of aid withdrawal should follow the geography of the aid economy. In host communities, where residents bear the costs of refugee presence but also receive jobs, services, and other benefits from refugee assistance, support for hosting should be relatively conditional on those material benefits. Aid withdrawal should therefore reduce support most in communities tied to the refugee aid economy and, within them, among households that lose aid-related benefits. Elsewhere in the country, where citizens are less directly exposed to either the costs or benefits of hosting, attitudes should be less contingent on refugee assistance and therefore less responsive to its withdrawal. This differs from a general economic-threat account of refugee attitudes. Whereas immigration research often emphasizes sociotropic concerns about the national economy (Hainmueller and Hopkins, 2014; Bansak, Hainmueller and Hangartner, 2023; Weber et al., 2025; Aviña et al., 2026), the mechanism here is egocentric and geographically concentrated: households that lose refugee-related benefits should become less supportive, producing localized rather than nationwide backlash.

This yields two hypotheses:

H1 (Community exposure): *Aid cuts reduce support for refugee hosting more in communities that depend on the refugee aid economy than in less exposed communities.*

H2 (Personal loss): *Within exposed communities, aid cuts reduce support for refugee hosting more among households that lose aid-related benefits than among households that do not.*

3 Context

3.1 Refugee Hosting in Kenya

I test these hypotheses in Kenya, a well-suited setting for three reasons. First, Kenya hosted approximately 800,000 refugees and was among the largest recipients of USAID assistance before the cuts (U.S. Department of State, 2025). Second, exposure to refugee assistance is geographically concentrated, with the vast majority of refugees, mostly from South Sudan and Somalia, living in the Kakuma and Dadaab camps (see Figure 1a; UNHCR, 2025). Camp-based refugees depend heavily on international assistance, given restrictions on movement outside the camps and limited employment opportunities (MacDonald and Lichtenheld, 2025).³

Third, existing research shows that refugees have brought substantial economic benefits to the area. Turkana is a semi-arid border county and one of Kenya’s poorest and most politically marginalized regions: 72% of residents live in poverty, compared with 68% of refugees and 37% nationwide (World Bank and UNHCR, 2021). The establishment of Kakuma camp in 1992 contributed to the growth of an urban center that is more developed than the rest of the county (International Finance Corporation, 2018; Alix-Garcia et al., 2018; Sanghi, Onder and Vemuru, 2016; World Bank and UNHCR, 2021).


Hosts in Kakuma benefit from the camp economy through three channels, which I describe using evidence from the panel survey (see Section 4) and World Bank and UNHCR (2021). First, hosts access refugee services, particularly health facilities and schools in the camp, shown in Figure 1b. 69% of Turkana hosts have no formal education (compared to 47% of refugees), and only 40% can read and write (78% for refugees). Hosts account for an estimated 10–15% of beneficiaries of camp healthcare services. As one respondent put it, “Since the refugees came, there has been help. Schools and hospitals were built” (Woman, 21).⁴

Second, hosts find employment with NGOs and with refugee households, particularly as casual domestic laborers. The Turkana hosts are “asked to fetch water, wash the


³13% of refugees live in urban areas, particularly Nairobi. This group is more self-sufficient and less dependent on aid.

⁴Refugees enabled hosts to “access basic needs like health from the UN” (Man, 31); “[we] get free treatments from the camp hospitals and free schooling for [our] children” (Man, 32).

Figure 1: Refugee-hosting areas and survey sample



(a) Refugee populations in Kenya



(b) Sample around Kakuma, by exposure status

Notes: Panel (a) shows refugee populations by hosting location and origin country. Panel (b) shows Kakuma boundaries, healthcare centers and schools, and respondent GPS locations by baseline-beneficiary status.

Table 1: Aid shocks and data collection

Date	Event	Description	Captured by	Trajectory
2023–2025	WFP reductions	General food basket cut in stages from 80% to 40% of minimum daily needs: 80→60% Jul 2023; 60→50% Feb 2024; 50→45% May 2024; 45→40% Feb 2025.	Natrep 2023 → 2026 (bundled)	Decline
Jan–Feb 2025	USAID cuts	USAID-funded NGOs reduced services and laid off staff across the camp.	Natrep 2023 → 2026 (bundled)	Decline
Jun 2025	Cash assistance suspended	Cash grants fully suspended; 40→30% food basket for all residents.	W1 → W2; Natrep	Decline
Jul 2025	DA announced	Vulnerability-tiered Differentiated Assistance model announced; sparked protests and delayed food distribution.	W1 → W2; Natrep	Trough
Aug 2025	DA takes effect	Tiered rations begin: Cat 1 at 40%, Cat 2 at 20%, Cat 3–4 receive nothing; cash still suspended; increased police presence.	W2 → W3; Natrep	Stabilization
Sep–Oct 2025	Cash resumption announced	Resumption of cash for Oct–Nov announced. Reduced insecurity.	W2 → W3; Natrep	Stabilization
Oct–Nov 2025	Rations increased, cash resumes	Rations raised (Cat 1 55%, Cat 2 35%, Cat 3 20%); cash resumes for Categories 1–3.	W3 → W4; Natrep	Partial reinstatement

dishes, do the laundry, and after doing the work [they] get something to go and eat at home with [their] kids” (Woman, 48).⁵ Third, hosts run businesses that depend on refugee demand, including shops and the sale of local resources. One respondent described: “Most of [us] take firewood to the camps and sell it to the refugees, and some take charcoal. After that, [we] are given food [by the refugees] and feed [our] kids” (Woman, 23).⁶ In May 2025, a large majority of panel respondents (71%) agreed that locals benefit from the camp.

3.2 Aid Cuts

Refugee assistance in Kenya was cut in three stages (see Table 1 for details).⁷ The third stage, captured by the panel, marked the trough of the cuts. Shortly after W1, WFP suspended the cash assistance refugees received alongside food rations. In August, it introduced a prioritization model that divided refugees into vulnerability tiers, leaving almost half of registered refugees without either rations or cash. Refugee leaders reported that many vulnerable households were misclassified as self-sufficient. The rollout triggered protests and delayed food distribution in late July, just prior to W2, contributing to acute hunger, malnutrition, and insecurity. By W3 in September, increased police presence had reduced insecurity, and partial aid reinstatement for some groups had been announced for October. Assistance nevertheless remained below previous levels: cash and rations resumed at reduced amounts, Category 4 refugees continued to receive no support, and NGO staffing did not recover.

Interviews with refugee leaders document the harm these cuts caused inside the camp, particularly in July 2025, establishing the magnitude of the shock that subsequently reached hosts. The central problem was widespread hunger, with many refugees eating “once per day”. To cope, they skipped meals, took on debt, resorted to crime, and sold their assets. One mother was forced to “sell the roof of the house . . . so that she can get food for her kids”; adolescent girls “[slept] with a man to get [sanitary] pads”; and “people have been cut with panga [knives] just because of a smart phone”. Schools and clinic access became more strained, the main issue being a lack of medication: “Now we have no drugs in the hospital. They’re just giving you painkillers. . . Malaria is the biggest killer”.

Refugees in Kakuma cannot move elsewhere in Kenya for work without formal permission, nor grow their own food, because “it’s arid, dry, [and] there’s no water.” Some returned to South Sudan despite the risk of conflict, reasoning that it is “better to die in your country with a bullet than hunger”. By late September, leaders noted the situation had improved, though not fully: “Comparing August and now, things are much better. August was really really bad. People had no food, there was a lot of tension and theft also. [Now I] feel like things are not good all time, but at least, you’re not hearing stories of someone starving and rushed to the hospital”. These accounts describe conditions inside the camp. In the rest of the paper, I examine how these refugee-aid cuts affected hosts and their attitudes toward refugees.

⁵Refugees “boosted the economy by introducing NGOs, creating job opportunities” (Man, 27).

⁶“Kakuma’s revenue has increased over time as compared to Lodwar due to a lot of businesses going on about and a lot of people duly paying their taxes” (Woman, 34).

⁷In fiscal year 2024, prior to the cuts, the U.S. provided 47.0% of WFP’s 2023–2027 contributions in Kenya and 45.8% of UNHCR’s contributions (WFP, 2025*b*; UNHCR, 2024). The U.S. disbursed \$116.7m to WFP, \$36.6m to UNHCR, \$4.8m to IRC, \$0.7m to LWF, and \$6.0m to other Kakuma organizations. Author calculated (U.S. Department of State, 2025; OCHA, 2025).

4 Data

I draw on three data sources to examine the effects of the aid cuts in Kenya on attitudes toward refugees. First, I fielded two original nationally representative phone surveys. The baseline was conducted in September–October 2023, with 3,326 respondents: a nationally representative sample of 2,432 and oversamples near Kakuma and Dadaab camps (441 in Garissa and 453 in Turkana).⁸ The follow-up was fielded in February 2026, after the major 2025 cuts, with a nationally representative sample of 2,042 and a Turkana oversample of 289.⁹ In both surveys, I construct post-stratification weights using entropy balancing, matching each sample to the contemporaneous Afrobarometer distribution of age, gender, education, and region (Hainmueller, 2012).¹⁰

Second, I fielded three waves of a face-to-face panel survey with Kenyan citizens in the host community adjacent to Kakuma camp.¹¹ The panel samples host residents only, not refugees. Although designed before the trajectory of aid was known, the panel tracks citizens through the main aid cut shock, from pre-trough conditions in May to the trough in July and minor improvements during stabilization in September. The main analysis uses 205 respondents present in all three waves. The Supplementary Materials reports parallel results using the repeated cross-section, refreshed to at least 300 citizens per wave. A fourth phone wave, fielded alongside the 2026 nationwide survey, captures longer-term change and partial aid reinstatement but has higher attrition and is reported in the Supplementary Materials. Figure 1b maps the camp, Kakuma town, panel respondents, and schools and health facilities.

Third, I conducted 52 semi-structured interviews with a panel of refugee leaders alongside the Kakuma survey waves. I use these interviews to document the timing and severity of the aid cuts inside the camp. Because hosts depend on the camp economy, this helps establish the shock underlying the employment and service losses for hosts. Refugee leaders are well placed to report this, since they are embedded in the distribution system and in regular contact with host residents, and the interviews are the source for the timing and content of the aid shocks reported in Table 1. Seventeen leaders remained in the panel across all three waves (18 in Wave 1 and 17 in Waves 2 and 3).¹² This evidence is informed by previous in-person interviews in Kakuma in 2022 and fieldwork across multiple years in Kenya.

⁸Because the 2023 survey randomized references to refugees generally or to specific groups, over-time comparisons use only the generic-refugees arm for *Nation Hosting* and *Local Hosting* ($N = 810$). See Supplementary Materials.

⁹A May 2026 follow-up corrected a technical issue in the *Nation Hosting* measure ($N = 1,776$).

¹⁰R9 (2021) and R10 (2024) for 2023 and 2026 surveys, respectively.

¹¹I use “Kakuma” as shorthand for Kakuma camp and neighboring Kalobeyei settlement. I was unable to study refugees in Dadaab due to security issues at the time of the study.

¹²Leaders were purposively sampled to vary by nationality, gender, and leadership role, including healthcare workers, local news broadcasters, leaders of refugee-led organizations, and employees of international organizations. Respondents were recruited through a local journalist and research assistant with strong networks in Kakuma. Details are in the Supplementary Materials.

I estimate all regressions by OLS, using a linear probability model for binary outcomes. Standard errors are HC1 for cross-sectional specifications and clustered by respondent for panel specifications. Specification details for the supplementary designs appear in the Supplementary Materials.

5 Comparing Hosting and Non-Hosting Regions

I begin by examining whether attitudes toward refugees in the country’s major hosting regions shifted after the cuts relative to the country as a whole. Using the two nationwide surveys, I examine change over a multi-year window, from 2023 to 2026, that spans the majority of the aid cuts, including most of the WFP ration reductions and the major 2025 cuts. The design exploits geographic variation in exposure to the aid economy: because community exposure is concentrated near the camps, comparing over-time change in the hosting regions with change nationwide isolates the differential effect of the cuts on the most exposed areas. Under H1, the hosting regions should exhibit more negative change than the rest of the country.

The outcome is the *Hosting Index*, an inverse-covariance-weighted index (Anderson, 2008) combining three items. *Nation Hosting* and *Local Hosting* are five-point scales measuring support for hosting refugees in Kenya and in the respondent’s district, respectively. *Not Repat.* is an indicator equal to one if the respondent selects any option other than repatriation (the most restrictive option, defined as returning refugees to their countries of origin) from among five refugee-policy alternatives: full rights, settlements, camps, repatriation, or third-country resettlement. All components are coded so that higher values indicate more supportive attitudes. The index is standardized to the 2023 national baseline (mean 0, SD 1), so coefficients represent standard-deviation changes relative to pre-cut attitudes. I also examine these items separately.

I estimate a pooled difference-in-differences comparing pre-cut (2023) and post-cut (2026) attitudes in the main refugee-hosting counties, Turkana and Garissa, with the rest of the country:¹³


$$Y_i = \alpha + \beta \text{Post}_i + \rho \text{Host}_i + \tau (\text{Post}_i \times \text{Host}_i) + X_i' \Gamma + \varepsilon_i, \quad (1)$$

where Host_i indicates a refugee-hosting county, and X_i controls for age, gender, and education. The coefficient β captures the adjusted change outside hosting counties, where community exposure is lower, while τ captures the differential change in hosting counties. H1 predicts $\tau < 0$ on the *Hosting Index*.

The nationwide difference-in-differences analysis finds strong support for H1. Table 2 shows that hosts in the refugee-hosting counties moved against refugees relative to

¹³Note that the sample size is larger for Turkana residents in the 2026 sample than Garissa residents, given that there was no Garissa oversample in this wave.

Figure 2: Hosting Index across Kenya, 2023–2026



Notes: Signed deviations in the *Hosting Index* from the substantive neutral point (all items at their midpoints) for 2023 and 2026, and the 2026–2023 change. Regions are entropy-balancing weighted. The Turkana overlay is unweighted and includes the oversamples. Red triangles mark Kakuma and Dadaab camps.

Table 2: Pooled difference-in-differences: nationwide refugee attitudes, 2023–2026

	Hosting Index	Nation Hosting	Local Hosting	Not Repat.
Post	0.114** (0.040)	-0.027 (0.066)	0.099 (0.072)	0.079*** (0.019)
Hosting county	0.326*** (0.063)	0.562*** (0.095)	0.381** (0.120)	0.109*** (0.028)
Post × Hosting county	-0.365*** (0.081)	-0.418* (0.196)	-0.659*** (0.155)	-0.099** (0.036)
Controls	Yes	Yes	Yes	Yes
Observations	3,401	2,841	3,381	3,357
R ²	0.040	0.023	0.034	0.033

Notes: OLS with HC1 standard errors; controls for age, gender, and education. All outcomes coded so that higher values indicate more pro-refugee attitudes. Full estimation results are reported in Supplementary Materials Section 2.3. *** $p < 0.001$, ** $p < 0.01$, * $p < 0.05$, † $p < 0.10$.

the rest of the country by roughly a third of a standard deviation ($\hat{\tau} = -0.365$, $p < 0.001$). The decline appears across all three components: it is largest for local hosting ($\hat{\tau} = -0.659$, $p < 0.001$) but also significant on nationwide hosting ($\hat{\tau} = -0.418$, $p < 0.05$) and not selecting repatriation as preferred policy ($\hat{\tau} = -0.099$, $p < 0.01$). Descriptively, support for local hosting in the camp counties fell from 61.3% to 48.4%, while nationwide support stayed stable (52.5% to 54.0%). Figure 2 shows regional shifts in support.

Because hosting regions began above the national average ($\hat{\rho} = 0.326$, $p < 0.001$), the shift represents convergence toward the national mean from a higher baseline rather than hosting regions becoming more hostile than the rest of the country. The coefficient for the non-hosting regions $\hat{\beta} = 0.114$ ($p < 0.01$) aligns with the prediction that where exposure is low and empathy is high, there is no backlash and even a slight increase in support, though the increase is substantively small and largely driven by the repatriation-choice measure.

6 Panel Evidence from Kakuma

6.1 Host-Community Effects of the Cuts


I next use the high-frequency panel to document how host households in Kakuma were affected during the most severe phase of the cuts, before examining changes in their hosting attitudes. As one host respondent put it, “life in the host community has started to become difficult. There is no food or work,” while a refugee leader observed that “the situation is becoming worse on both sides.” Figure 3 shows declines in services, earnings, employment, and security across the panel. Access to camp healthcare fell from 42.4% at W1 to 26.8% at W3, while the share feeling unsafe at night rose from 58.1 to 71.6%. Median monthly household earnings fell from KSh 9,000 to KSh 5,000 and did not recover, while the share earning below KSh 5,000 rose from 33.9 to 47.3%. Household-head unemployment also rose at W2, although it largely returned to baseline by W3.

Conditions improved somewhat between W2 and W3, but most lost benefits had not returned. Among households that lost camp-area education, healthcare, refugee employment, or NGO employment after W1, only 8% had regained access in that domain by W3, and earnings remained depressed. Crime victimization did decline between W2 and W3, alongside the recovery in employment, indicating some stabilization after the protests and insecurity surrounding W2.

6.2 Within-Person Changes in Hosting Attitudes

I next examine whether these changes to household outcomes affect attitudes toward local refugee hosting. The analysis uses within-person variation across three closely

Figure 3: Host livelihood outcomes across the Kakuma panel



Notes: Share of Kakuma panel households reporting each outcome by wave, with 95% confidence intervals from linear probability models and standard errors clustered by respondent. Services, employment, and crime refer to the previous month. Severe food insecurity indicates missing meals on five to seven days in the previous week. Full estimates and wave-to-wave tests appear in Supplementary Materials Section 3.2.

spaced panel waves: W1 captures conditions before the trough, W2 the most severe phase of the contraction, and W3 continued cuts under somewhat more stable conditions. Under H1, attitudes in the hosting region should worsen once aid is cut. Under H2, the decline should be concentrated among those who personally lose aid-related benefits. The high-frequency nature of the panel lends confidence that attitudinal shifts track the specific aid changes documented in Section 3.2 rather than slower-moving confounders.

I focus on two outcomes specific to local hosting, both of which capture how residents feel about continuing to host refugees in their area. *Oppose Camp Closure* measures disagreement with “the Kenyan government closing Kakuma camp and moving the refugees in Kakuma back to their countries,” recorded on a five-point scale and coded so that higher values indicate pro-refugee attitudes. *Local Hosting* is the same district-level hosting item described above. I also report *Nation Hosting* for comparison with

support for hosting in the country more broadly.

I estimate an event-study specification on the balanced panel,

$$Y_{it} = \mu + \nu_i + \delta_2 D_{2t} + \delta_3 D_{3t} + \varepsilon_{it}, \quad (2)$$

with respondent fixed effects ν_i , wave indicators D_{2t} and D_{3t} (W1 omitted). The coefficient δ_2 is the cumulative change from baseline to the trough of the contraction, and the linear contrast $\delta_3 - \delta_2$ measures the change from the trough to the stabilization period. H1 predicts $\delta_2 < 0$ for both *Local Hosting* and *Oppose Camp Closure*.

Where H1 concerns whether community exposure is associated with backlash, H2 asks whether the backlash within high-exposure communities is driven by the households that personally lose aid-related benefits. I test this with an intent-to-treat difference-in-differences design comparing households vulnerable to personal loss with those that are not. I define the treated group by baseline attributes: a household is a "baseline-beneficiary" if, at W1, the respondent accessed camp healthcare or education services or was employed by an NGO or a refugee in the past 30 days. This applies to 127 of 205 respondents (62.0%).¹⁴ Because status is fixed at baseline, before the severe cuts, the design avoids post-treatment bias and is likely conservative, since not every baseline beneficiary lost benefits. I examine the same outcomes as above.

I interact a baseline-beneficiary indicator with wave dummies on the balanced panel, with respondent fixed effects:

$$Y_{it} = \mu + \nu_i + \delta_2 D_{2t} + \delta_3 D_{3t} + \tau_2 (B_i \times D_{2t}) + \tau_3 (B_i \times D_{3t}) + \varepsilon_{it}, \quad (3)$$

where B_i equals one if the household received camp healthcare, education, or employment in the 30 days before W1. The coefficients τ_2 and τ_3 capture the differential change among baseline beneficiaries at the trough and stabilization waves, and H2 predicts $\tau_2 < 0$ and $\tau_3 < 0$.

Selection into beneficiary status is non-random: households receiving aid-related benefits before the cuts were more supportive of refugees at baseline than those that were not. Respondent fixed effects absorb this baseline gap, but cannot rule out that the two groups were already on different trajectories, so I read these estimates as quasi-causal and triangulate with two further designs. The first defines treatment as households that actually lost a benefit, comparing those that lost access between W2 and W3 with those that never lost access, using the W1–W2 window to check for parallel pre-trends. The second pools all loss events, estimating how attitudes change when a household switches into or out of benefit access between any adjacent waves. See Supplementary Materials for details.

¹⁴Camp healthcare 42.4%, camp education 20.0%, NGO employment 19.5%, refugee employment 13.2%. Categories overlap.

Table 3: Within-Kakuma hosting attitudes over time and by household loss

	Oppose Camp Closure	Local Hosting	Nation Hosting
<i>Panel A: Over-time change (event study)</i>			
Wave 2 vs Wave 1	-0.260 [†] (0.151)	-0.010 (0.128)	-0.166 (0.136)
Wave 3 vs Wave 2	-0.039 (0.147)	-0.059 (0.167)	0.441** (0.135)
Respondent-wave obs.	612	613	614
Respondents	204	205	205
<i>Panel B: Intent-to-treat, baseline beneficiary status</i>			
Wave 2	0.179 (0.251)	-0.051 (0.219)	-0.103 (0.237)
Wave 3	0.141 (0.266)	0.332 (0.288)	0.645* (0.270)
Baseline beneficiary (W1)	0.953*** (0.258)	0.926*** (0.247)	0.735** (0.252)
Baseline beneficiary $\times$ Wave 2	-0.697* (0.312)	0.067 (0.269)	-0.102 (0.289)
Baseline beneficiary $\times$ Wave 3	-0.699* (0.336)	-0.669 [†] (0.352)	-0.598 [†] (0.334)
Respondent-wave obs.	613	613	614
Respondents	205	205	205

Notes: Panel A: OLS, respondent fixed effects, standard errors clustered by respondent. Panel B: pooled OLS, standard errors clustered by respondent so that the time-invariant baseline-beneficiary difference is identified ($N = 127$ beneficiary, $N = 78$ non-beneficiary). No other covariates. Higher values = more pro-refugee attitudes. *** $p < 0.001$, ** $p < 0.01$, * $p < 0.05$, [†] $p < 0.10$.

Table 3 (Panel A) shows some support for H1. There is a decline in opposition to camp closure from W1 to W2 ($\hat{\delta}_2 = -0.260$, $p < 0.10$), consistent with backlash. This is only weakly significant, though this may be due to the small sample. There is no significant decrease in support for local hosting over the aid cuts period. Neither measure changes significantly between W2 and W3, meaning that elevated support for closure persisted through to W3.

However, when we examine those who were most vulnerable to economic loss within the hosting region, we see significant declines in support across all hosting measures. Panel B of Table 3 supports H2. Beneficiary households were substantially more supportive before the cuts but became less so over time. On the camp-closure outcome, beneficiaries shifted toward closure at both W2 ($\hat{\tau}_2 = -0.697$, $p < 0.05$) and W3 ($\hat{\tau}_3 = -0.699$, $p < 0.05$). By W3, similar declines appeared for *Local Hosting* ($\hat{\tau}_3 = -0.669$, $p < 0.10$) and *Nation Hosting* ($\hat{\tau}_3 = -0.598$, $p < 0.10$). These shifts nearly eliminated the pre-cut gap in support for local hosting between beneficiaries and non-beneficiaries. By contrast, attitudes among non-beneficiaries changed little across waves. Descriptively, support for closing Kakuma camp rose from 29.1% at W1 to 45.2% at W3 among baseline beneficiaries, while remaining stable among non-beneficiaries (53.8% to 48.7%). The two additional treatment specifications in the Supplementary Materials replicate these findings. The backlash documented under H1 was therefore concentrated among households that personally lost from the cuts, while attitudes among non-beneficiaries remained stable.


7 Experimental Evidence

To test causally whether aid cuts reduce support for refugees, I embed a factorial vignette experiment in W3 of the Kakuma panel. The experiment asks whether residents of the hosting region withdraw support when the aid attached to hosting changes. Respondents rate their support for Kenya hosting refugees (the *Nation Hosting* measure) under scenarios that randomly vary four attributes: whether international organizations provide 100, 40, or 0% of food rations and cash support to refugees (*Aid*); whether healthcare and schools are shared with the host community or reserved for refugees only (*Services*); whether international organizations hire local residents (*Hiring*); and whether conflict is ongoing or peace has been restored in refugees' country of origin (*Conflict*). The attribute of primary interest for H1 is *Aid*: support for hosting should fall as aid drops from 100% to lower levels.

The $3 \times 2 \times 2 \times 2$ design yields 24 possible profiles. To limit respondent fatigue, each respondent rates six profiles drawn at random. I estimate average marginal component effects (AMCEs) with respondent fixed effects. Respondents are asked a comprehension check concerning the level of aid in the final scenario they rate.

The experiment finds support for H1. Figure 4 reports the AMCEs for the vignette

Figure 4: Vignette experiment



Notes: AMCEs from Wave 3 factorial vignette (Kakuma, $N = 299$ respondents, 1,883 profiles). Points show mean *Nation Hosting* at each attribute level; dashed line marks reference. Linear model with respondent fixed effects, SEs clustered by respondent. The full estimation results are reported in Supplementary Materials Section 4.2. *** $p < 0.001$, ** $p < 0.01$, * $p < 0.05$, † $p < 0.10$.

experiment and shows that all three reductions in aid for hosts move attitudes in the predicted direction: cutting aid to zero lowers hosting support by 0.19 points ($p < 0.05$), removing local hiring by 0.20 ($p < 0.01$), and restricting services to refugees by 0.18 ($p < 0.05$). Reducing aid to 40% produces a smaller decline of 0.14 that is weakly significant ($p < 0.10$).


I complement the experiment with a descriptive item, fielded in W1 and W2, asking respondents whether they would support Kenya hosting refugees if international aid were withdrawn entirely. The pattern matches the experimental results. In W2, 50.2% of respondents supported hosting refugees in Kenya, but only 29.4% did so under a full withdrawal of aid, a within-respondent decline of 0.84 points on the five-point scale ($p < 0.001$; Supplementary Materials Section 4.4).

8 Qualitative Data

The designs above establish that support for hosting fell where exposure to the aid economy was high, and that within the hosting region the decline was concentrated among households that personally lost benefits. To understand the mechanisms underlying these changes, as well as potentially why support stayed stable in the rest of the country, I employ three open-ended questions. In the Kakuma panel, I ask respondents why they support or oppose closing the camps, following the *Oppose Camp Closure* question (W3), and what changes they have seen from the cuts (W1–W3). The third question, asked nationwide in 2023, is why respondents support or oppose hosting refugees.¹⁵ Responses may receive multiple codes, so shares sum to more than

¹⁵The national question was fielded before the cuts, so it characterizes the pre-cut basis of national support rather than change in it.

Figure 5: Open-ended responses



Notes: Share of respondents mentioning each theme. Panels A and B: W3 Kakuma respondents opposing ($N = 163$) and supporting ($N = 124$) camp closure. Panel C: pooled W1–W3 Kakuma ($N = 615$ respondent-waves). Panel D: 2023 national sample, control arm. Coding details are in Supplementary Materials Section 1.4.2.

one.¹⁶

The open-ended responses in Figure 5 show three patterns. First, responses to the camp-closure question emphasized economic benefits and insecurity. Among those who opposed closure, 63.8% cited the camps’ economic benefits, suggesting that support for local hosting depends substantially on its material benefits. One respondent opposed closure because “[the] camp helps a lot of people like those who sell charcoal and firewood, and a lot [of] others benefit from the casual jobs over there . . . if closed then crime rates will rise tremendously” (Man, 38). Others drew the opposite conclusion:

¹⁶I create a set of categories for these questions by examining 10% of responses and then use Claude to code the remaining responses into these categories, which I subsequently check.

among those who supported closure, insecurity was the most common explanation, cited in 58.1% of responses. These respondents linked insecurity directly to the cuts, describing refugees who “have no food and most of them have become thieves” (Man, 25) and a “demonstration” in which they “wanted to burn down the camp” (Man, 19). On economic losses, one respondent said “there’s no more jobs, no relief food coming, businesses [have] deteriorated, so what’s the use of still keeping them here” (Woman, 35).

Second, Kakuma residents described the cuts mainly in terms of their effects on refugees. Refugee suffering was the most commonly mentioned change, cited in 34.1% of respondent-waves. However, a large proportion also spoke about crime, protests, and violence, as well as a lack of food for both themselves and the refugees.

Third, national responses placed greater emphasis on humanitarian considerations. More than 40% cited humanitarian reasons, more than any other category and more than the economic categories combined. Material benefits therefore appear less central outside the main hosting regions, which may help explain why support declined in hosting regions but not elsewhere in the nationwide difference-in-differences result.

9 Humanitarian Concerns

The preregistered analyses were not designed to test humanitarian concern, but three findings suggest that it may shape attitudes alongside material interest: nationwide support remained stable, open-ended responses frequently invoked refugee suffering, and ongoing conflict in refugees’ countries of origin produced the largest experimental effect (Figure 4). I therefore explore this possibility through two additional analyses.

First, I examine whether the experimental effect of aid withdrawal varies with empathy for refugees. The decline in support under the 40% aid condition is concentrated among low-empathy respondents: the interaction between empathy and the 40% condition is positive and significant (0.230, $p < 0.01$; Supplementary Materials Section 4.5). Descriptively, support among low-empathy respondents falls from 40.9% under full aid to 32.5% under 40% aid, whereas support among high-empathy respondents remains near 40% across aid conditions.

Second, I find that although Kakuma residents became more supportive of closing the camp, they did not become less supportive of refugees more generally. Willingness to host refugees in the country more broadly, trust in refugees, and support for refugee rights remained stable (Table 3; Supplementary Materials).

These patterns align with research showing that humanitarian need, empathy, and perspective-taking can increase support for refugees (Bansak, Hainmueller and Hangartner, 2016, 2023; Newman et al., 2015; Hartman and Morse, 2020; Hartman, Morse and

Weber, 2021; Adida, Lo and Platas, 2018; Williamson et al., 2021).

10 Robustness

The Supplementary Materials reports the full set of robustness checks. For the nationwide difference-in-differences, samples are nationally representative and comparable to each other. Because only two counties are treated, I use a wild cluster bootstrap (Cameron, Gelbach and Miller, 2008). The estimates are stable across alternative index constructions, outcome codings, estimators, weighting schemes, and county leave-one-out samples. I show that the nationwide results are unlikely to be related to the Shirika Plan. I also report additional outcomes, including an Afrobarometer refugee-numbers item fielded in 2024 and 2026, and a longer-run follow-up fielded alongside the February 2026 national survey.

For the panel, no Wave 1 observable predicts attrition, and inverse-probability-of-retention weights leave the estimates essentially unchanged. The results replicate on the full repeated cross-section of at least 300 respondents per wave. Two further personal-loss specifications, a cohort difference-in-differences with a Wave 2 placebo and a switching design pooling loss events by domain, corroborate the intent-to-treat result, which also holds under enumerator fixed effects. For the vignette experiment, treatment assignment is random and 97.7% of respondents pass the comprehension check, with results unchanged when those who fail are excluded.

Given the small sample, I report minimum detectable effects for the main tests. Because the quantities of interest are over-time changes rather than levels, social desirability pressures present at every wave difference out. And because both the hosting counties and the baseline-beneficiary households begin with more favorable attitudes, I also address the possibility that the estimated declines reflect reversion to the mean.

11 Conclusion

Although governments decide nationwide refugee policy, the consequences of hosting are often concentrated in a small number of communities that share land, services, markets, and security conditions with refugee populations. This paper has argued that cooperation with the refugee regime therefore depends partly on sustaining support in these communities. Refugee aid can help do so by bringing jobs, services, infrastructure, and other benefits to places that bear many of the costs of hosting. The sudden contraction of humanitarian assistance in 2025 allows us to examine what happens when that support is withdrawn.

Drawing on two nationwide surveys, a high-frequency panel in a refugee-hosting region, a vignette experiment, and interviews with refugee leaders in Kenya, I report three

main findings. First, support for refugee hosting in host communities depends in part on the material benefits accompanying refugee aid. Second, the aid cuts generated a localized backlash against refugee hosting. Support declined in hosting regions, driven by households that lost aid-related jobs or services, and qualitative evidence links this decline to economic losses and growing insecurity. Third, this backlash did not extend nationwide, where support remained stable. Exploratory evidence suggests that humanitarian concern may limit backlash.

These findings show why national public opinion can provide a misleading picture of the political foundations of refugee cooperation. The communities that host refugees are a small share of the national population, so declining local support may not significantly affect aggregate attitudes. Yet local opposition can matter disproportionately for policy. Grievances in hosting areas can generate protests, insecurity, conflict, and demands for camp closure or relocation, creating pressure on governments to restrict refugee protection even when broader public opinion remains stable.

The findings also speak to the durability of international cooperation with the refugee regime. Existing accounts emphasize the bargain between host states that provide asylum and donor states that finance refugee protection. The results here suggest that this interstate arrangement also rests on a subnational political foundation. Where refugees are geographically concentrated, cooperation may depend on whether host communities receive sufficient benefits to offset the local costs. Because many of those benefits are externally financed, donor retrenchment can weaken cooperation.

The dynamics documented here are unlikely to be unique to Kenya, given the severe cuts to refugee assistance across LMICs, but should be most pronounced where refugee hosting, foreign aid, and its associated benefits are geographically concentrated. This is particularly true in camp-based settings, where surrounding communities often share in the jobs, services, and economic activity generated by aid. Roughly one in five refugees worldwide, about eight million people, lives in a camp, particularly in sub-Saharan Africa but also in countries such as Jordan, Thailand, and Bangladesh (Anti et al., 2025; UNHCR, 2026*b*). The effects should be less pronounced where refugees are more dispersed or where host communities are less economically tied to humanitarian assistance. Cross-national research is needed to establish under what conditions donor withdrawal weakens the local foundations of refugee cooperation.

The central takeaway is that aid cuts have political consequences beyond their direct welfare effects. In Kenya, cuts to food, services, and employment spilled over to surrounding communities through falling household incomes, reduced access to schools and hospitals, rising insecurity, and protests. These changes weakened support for continued local refugee hosting among the households and communities most exposed to the aid economy, even as national attitudes remained stable. Aid therefore does more than finance refugee protection: it can help sustain cooperation in the places where refugee protection is actually implemented. When donors withdraw that support, they

risk weakening the local political foundations on which the international refugee regime ultimately depends.

References

- Adida, Claire L., Adeline Lo and Melina R. Platas. 2018. “Perspective Taking Can Promote Short-Term Inclusionary Behavior toward Syrian Refugees.” *Proceedings of the National Academy of Sciences* 115(38):9521–9526.
- Alix-Garcia, Jennifer, Sarah Walker, Anne Bartlett, Harun Onder and Apurva Sanghi. 2018. “Do refugee camps help or hurt hosts? The case of Kakuma, Kenya.” *Journal of Development Economics* 130:66–83.
- Anderson, Michael L. 2008. “Multiple Inference and Gender Differences in the Effects of Early Intervention: A Reevaluation of the Abecedarian, Perry Preschool, and Early Training Projects.” *Journal of the American Statistical Association* 103(484):1481–1495.
- Anti, Sebastian, Colette Salemi, Jonathan Rigberg, Karishma Silva and Johannes Hoogeveen. 2025. “A New Expanded Dataset to Study Refugee Camps in Sub-Saharan Africa 1999–2024.” Working paper.
URL: https://colettesalemi.ca/files/Anti_etal_2025.pdf
- Arar, Rawan. 2017. “The New Grand Compromise: How Syrian Refugees Changed the Stakes in the Global Refugee Assistance Regime.” *Middle East Law and Governance* 9(3):298–312.
- Ariba, Caroline. 2026. “Less Drugs, Less Care, Less Food: How Aid Cuts Have Hit Uganda’s 2m Refugees.” *The Guardian*. 9 July.
URL: <https://www.theguardian.com/global-development/2026/jul/09/how-aid-cuts-have-hit-uganda-2m-refugees>
- Associated Press. 2026. “Food Assistance Slashed for Hundreds of Thousands of Rohingya Refugees Trapped in Bangladesh Camps.” Associated Press, via NPR. 1 April.
URL: <https://www.npr.org/2026/04/01/nx-s1-5769798/food-assistance-slashed-rohingya-refugees-bangladesh-camps>
- Aviña, Marco M., Taeku Lee, Mashail Malik, Reed Rasband, Marcel F. Roman and Priyanka Sethy. 2026. “Which Immigrants Do Citizens Prefer? A Meta-Reanalysis of 100 Conjoint Experiments.” *Science Advances* 12(22).
- Bansak, Kirk, Jens Hainmueller and Dominik Hangartner. 2016. “How economic, humanitarian, and religious concerns shape European attitudes toward asylum seekers.” *Science* 354(6309):217–222.
- Bansak, Kirk, Jens Hainmueller and Dominik Hangartner. 2023. “Europeans’ support for refugees of varying background is stable over time.” *Nature* 620(7975):849–854.
- Baseler, Travis, Thomas Ginn, Robert Hakiza, Helidah Ogude-Chambert and Olivia

- Woldemikael. 2025. “Can redistribution change policy views? Aid and attitudes toward refugees.” *Journal of Political Economy* 133(9):2874–2917.
- Betts, Alexander. 2003. “Public Goods Theory and the Provision of Refugee Protection: The Role of the Joint-Product Model in Burden-Sharing Theory.” *Journal of Refugee Studies* 16(3):274–296.
- Betts, Alexander. 2009. *Protection by Persuasion: International Cooperation in the Refugee Regime*. Ithaca, NY: Cornell University Press.
- Broz, J. Lawrence, Jeffrey Frieden and Stephen Weymouth. 2021. “Populism in Place: The Economic Geography of the Globalization Backlash.” *International Organization* 75(2):464–494.
- Cameron, A. Colin, Jonah B. Gelbach and Douglas L. Miller. 2008. “Bootstrap-Based Improvements for Inference with Clustered Errors.” *The Review of Economics and Statistics* 90(3):414–427.
- Ceritoglu, Evren, H Burcu Gurcihan Yunculer, Huzeyfe Torun and Semih Tumen. 2017. “The impact of Syrian refugees on natives’ labor market outcomes in Turkey: evidence from a quasi-experimental design.” *IZA Journal of Labor Policy* 6:1–28.
- Cuéllar, Mariano-Florentino. 2006. “Refugee Security and the Organizational Logic of Legal Mandates.” *Georgetown Journal of International Law* 37:583–723.
- Demir, Cansu, Anna Maria Mayda and Jean-François Maystadt. 2025. “The Political Impact of Refugees in Africa.” Working paper.
URL: <https://www.jointdatacenter.org/wp-content/uploads/2026/06/Jean-Francois-Maystadt.pdf>
- Farge, Emma. 2026. “Over 1 Million Sudanese Refugees in Chad Face Drastic Aid Cuts, UN Says.” Reuters. 9 April.
URL: <https://www.reuters.com/world/africa/over-1-million-sudanese-refugees-chad-face-drastic-aid-cuts-un-says-2026-04-09/>
- Ferreira da Silva, Andrea, Rodrigo Volmir Rezende Anderle, Gonzalo Barreix Sibils, Lucas de Oliveira Ferreira de Sales, Daiana Pena, Caterina Monti et al. 2026. “Impact of Two Decades of Humanitarian and Development Assistance and the Projected Mortality Consequences of Current Defunding to 2030: Retrospective Evaluation and Forecasting Analysis.” *The Lancet Global Health* 14(5):e690–e701.
- Ghosn, Faten, Alex Braithwaite and Tiffany S Chu. 2019. “Violence, displacement, contact, and attitudes toward hosting refugees.” *Journal of Peace Research* 56(1):118–133.
- Grossman, Guy, Yang-Yang Zhou and Shelby Carvalho. 2025. “Disrupted Aid, Displaced Lives: Unraveling the Impact of Refugee Funding Cuts through a Regression Discontinuity Study in Uganda.” Working paper.

- URL:** <https://pdri-devlab.upenn.edu/project/disrupted-aid-displaced-lives-unraveling-the-impact-of-refugee-funding-cuts-through>
- Hainmueller, Jens. 2012. “Entropy balancing for causal effects: A multivariate reweighting method to produce balanced samples in observational studies.” *Political Analysis* 20(1):25–46.
- Hainmueller, Jens and Daniel J Hopkins. 2014. “Public attitudes toward immigration.” *Annual Review of Political Science* 17(1):225–249.
- Hartman, Alexandra C and Benjamin S Morse. 2020. “Violence, empathy and altruism: Evidence from the Ivorian refugee crisis in Liberia.” *British Journal of Political Science* 50(2):731–755.
- Hartman, Alexandra C, Benjamin S Morse and Sigrid Weber. 2021. “Violence, displacement, and support for internally displaced persons: Evidence from Syria.” *Journal of Conflict Resolution* 65(10):1791–1819.
- International Finance Corporation. 2018. Kakuma as a Marketplace: A Consumer and Market Study of a Refugee Camp and Town in Northwest Kenya. Technical report International Finance Corporation Washington, DC: . License: All Rights Reserved.
URL: <https://openknowledge.worldbank.org/handle/10986/34645>
- Jordan Times. 2026. “UNHCR Jordan Reports \$202 Million Funding Gap as Refugee Operations Continue.” Jordan Times. 23 July.
URL: <https://jordantimes.com/news/local/unhcr-jordan-reports-202-million-funding-gap>
- Kadigo, Mark Marvin and Jean-Francois Maystadt. 2023. “How to cope with a refugee population? Evidence from Uganda.” *World Development* 169:106293.
- Kenny, Charles and Justin Sandefur. 2025. Update on Lives Lost from USAID Cuts. CGD blog post Center for Global Development.
URL: <https://www.cgdev.org/blog/update-lives-lost-usaid-cuts>
- MacDonald, Mae and Adam Lichtenheld. 2025. “Support for Refugee Integration in a Major Refugee-Hosting Country: Evidence from Kenya.” Working paper.
URL: https://osf.io/preprints/socarxiv/92qy5_v3
- Newman, Benjamin J., Todd K. Hartman, Patrick L. Lown and Stanley Feldman. 2015. “Easing the Heavy Hand: Humanitarian Concern, Empathy, and Opinion on Immigration.” *British Journal of Political Science* 45(3):583–607.
- OCHA. 2025. “Financial Tracking Service (FTS).”. Accessed: 20 March 2025.
URL: <https://fts.unocha.org>
- OECD. 2026. “A Historic Decline in Foreign Aid: Preliminary 2025 ODA Data.” <https://www.oecd.org/en/data/insights/data-explainers/2026/04/a-historic-decline-in-foreign-aid-preliminary-2025-oda-data.html>. Accessed 23 May 2026.

Pettigrew, Thomas F. and Linda R. Tropp. 2006. “A Meta-Analytic Test of Intergroup Contact Theory.” *Journal of Personality and Social Psychology* 90(5):751–783.

Putnam, Robert D. 1988. “Diplomacy and Domestic Politics: The Logic of Two-Level Games.” *International Organization* 42(3):427–460.

Regan, Helen, Ivan Watson, Kocha Olarn, Eugene Sein and Mark Phillips. 2025. “No Doctors for Sick Children. This Is the Reality of Trump’s Aid Freeze in Remote Northern Thailand.” CNN. 11 February.

URL: <https://www.cnn.com/2025/02/11/asia/thailand-myanmar-refugees-us-aid-freeze-trump-intl-hnk-dst>

Ruggie, John Gerard. 1982. “International Regimes, Transactions, and Change: Embedded Liberalism in the Postwar Economic Order.” *International Organization* 36(2):379–415.

Salehyan, Idean and Kristian Skrede Gleditsch. 2006. “Refugees and the spread of civil war.” *International organization* 60(2):335–366.

Sanghi, Apurva, Harun Onder and Varalakshmi Vemuru. 2016. “Yes” in My Backyard? The Economics of Refugees and Their Social Dynamics in Kakuma, Kenya. Technical report World Bank Washington, DC: .

URL: <https://documents1.worldbank.org/curated/en/308011482417763778/pdf/111303-WP-Kakuma-Report-Yes-in-my-backyard-December-2016-PUPdf>

Sterck, Olivier and Vittorio Bruni. 2025. “The Welfare and General-Equilibrium Impacts of Aid Cuts.” Working Paper.

URL: <https://ssrn.com/abstract=5270846>

Suhrke, Astri. 1998. “Burden-sharing during Refugee Emergencies: The Logic of Collective versus National Action.” *Journal of Refugee Studies* 11(4):396–415.

UNHCR. 2024. “Kenya Funding Update.” <https://reporting.unhcr.org/kenya-funding-update>. Accessed: 2025-03-21.

UNHCR. 2025. “Kenya Operational Data Portal.”.

URL: <https://data.unhcr.org/en/country/KEN/796>

UNHCR. 2026a. Global Report 2025. Technical report United Nations High Commissioner for Refugees Geneva: . Accessed July 2026.

URL: <https://www.unhcr.org/publications/global-report-2025>

UNHCR. 2026b. “Operational Data Portal: Refugee Situations.” <https://data.unhcr.org>. Accessed July 2026.

U.S. Department of State. 2025. “Foreign Assistance Data.”. Accessed: 20 March 2025.

URL: <https://www.foreignassistance.gov>

- Verme, Paolo and Kirsten Schuettler. 2021. "The impact of forced displacement on host communities: A review of the empirical literature in economics." *Journal of Development Economics* 150:102606.
- Weber, Sigrid, Nik Stoop, Peter Van der Windt and Haoyu Zhai. 2025. "A meta-analysis of attitudes towards migrants and displaced persons." *British Journal of Political Science* 55:e168.
- WFP. 2025a. "Food Security Impact of Reduction in WFP Funding." World Food Programme.
URL: <https://www.wfp.org/publications/food-security-impact-reduction-wfp-funding>
- WFP. 2025b. "Kenya Country Portfolio Needs (2023–2027)." World Food Programme, Operation KE02 Resourcing Update. Accessed March 20, 2025.
URL: https://one.wfp.org/operations/current_operations/ResUpdates/KE02.pdf
- Williamson, Scott, Claire L. Adida, Adeline Lo, Melina R. Platas, Lauren Prather and Seth H. Werfel. 2021. "Family Matters: How Immigrant Histories Can Promote Inclusion." *American Political Science Review* 115(2):686–693.
- World Bank and UNHCR. 2021. Understanding the Socioeconomic Conditions of Refugees in Kenya: Volume B: Kakuma Camp. Results from the 2019 Kakuma Socioeconomic Survey. Technical report World Bank and UNHCR. Accessed November 14, 2024.
URL: <https://documents1.worldbank.org/curated/en/443431613628051180/pdf/Understanding-the-Socio-Economic-Conditions-of-Refugees-in-Kenya-Volume-B-Kakuma-Camp-Results-from-the-2019-Kakuma-Socioeconomic-Survey.pdf>
- Zhou, Yang-Yang, Guy Grossman and Shuning Ge. 2023. "Inclusive refugee-hosting can improve local development and prevent public backlash." *World Development* 166:106203.
- Zhou, Yang-Yang, Naijia Liu, Shuning Ge and Guy Grossman. 2025. "Liberalizing Refugee Hosting Policies without Losing the Vote." OSF Preprints, Version 4.
URL: <https://doi.org/10.31219/osf.io/94tpy>

Aid Cuts and the Politics of Refugee Hosting

Supplementary Materials

Contents

1	Study design and measurement	3
1.1	Enumeration and recruitment	3
1.1.1	Nationwide	3
1.1.2	Panel	3
1.1.3	Refugee-leader interviews	3
1.2	Ethics	4
1.3	Sample	5
1.3.1	Nationwide surveys	5
1.3.2	Kakuma	6
1.4	Survey instruments	7
1.4.1	Nationwide outcome items	8
1.4.2	Open-ended questions and coding	9
1.4.3	Vignette experiment (W3)	9
1.5	Estimation equations	10
1.5.1	Nationwide difference-in-differences (H1)	10
1.5.2	Kakuma panel event study (H1)	10
1.5.3	Personal loss (H2)	11
1.5.4	Vignette experiment (H1)	11
1.6	Hosting Index construction	12
2	Nationwide results and robustness	12
2.1	Summary statistics	12
2.2	Weighted difference-in-differences analysis	13
2.3	Full estimation results for the main difference-in-differences	13
2.4	Covariate-by-period interactions	14
2.5	Wild cluster bootstrap	15
2.6	County leave-one-out analysis	16
2.7	Additional dependent variables	17
2.7.1	Hosting index robustness	17
2.7.2	Afrobarometer refugee-numbers measure	17
2.7.3	Integration measures	18
2.8	Alternative estimators	19
3	Kakuma panel results and robustness	19
3.1	Summary statistics	19
3.2	Over-time welfare tests	20
3.3	Repeated cross-section	21
3.4	Fourth wave: longer-term changes	21
3.5	Integration and trust measures	23
3.6	Personal-loss difference-in-differences	23
3.6.1	Comparing baseline-beneficiaries with non-baseline-beneficiaries	23
3.6.2	Alternative personal loss measures	24

4	Survey experiment results and robustness	26
4.1	Treatment arm balance	26
4.2	Full AMCE estimation results	26
4.3	Comprehension check	26
4.4	Hypothetical full-cut scenario	27
4.5	Humanitarian concern and empathy	29
4.5.1	Empathy index construction and robustness	29
4.5.2	Empathy marginal means	30
5	Additional analyses and study documentation	31
5.1	Enumerator fixed effects	31
5.2	Minimum detectable effects (MDE)	32
5.3	Shirika Plan	33
5.4	Reversion to the mean	33
5.5	Deviations from the pre-analysis plan	34
	References	37

1 Study design and measurement

1.1 Enumeration and recruitment

1.1.1 Nationwide

I worked with the Nairobi-based survey firm, TIFA Research, to conduct this study. TIFA maintains a database of previously profiled citizens, from which it recruited for this project. Respondents were drawn by stratified random sampling on gender, age, and location so that the sample matched the demographic distribution of the 2019 Kenya census. The sample is distributed across counties in proportion to the adult population using probability proportional to size, with the 2019 Kenya Population and Housing Census as the reference. Citizens were eligible if they were Kenyan voters aged 18 to 65.

Prior to data collection, I conducted extensive in-person piloting and training with enumerators. TIFA implemented rigorous quality control checks, including a quality control team listening to every interview to check for enumerator miscoding and to ensure that the qualitative open-ended responses had been captured correctly. If an interview was incorrectly administered, it was removed and not counted in the total survey number. For the qualitative open-ends, enumerators asked these questions in Swahili or English (or Somali in the 2023 survey, because of the Garissa oversample) depending on the preferences of the respondent, and then typed the response into TIFA’s survey software in that language. TIFA’s quality control team checked the correctness and quality of these responses against their recordings before employing translators to translate the responses into English for analysis. Ethics approvals, consent procedures, incentives, and data protections are described in Section 1.2.

In the 2023 baseline, the *Nation Hosting* and *Local Hosting* measures were embedded in a survey experiment fielded for a separate study, in which respondents were randomized across arms that varied refugee nationality. In the current study, I use only the generic-refugees control arm for these outcomes, which reduces the 2023 sample from the full nationally representative sample of 2510 to the 816 respondents in the control arm. A technical error in the February 2026 survey affected the *Nation Hosting* measure only. The fault for this error is mine, not TIFA’s. Because of this, I fielded a recontact survey in May 2026, which asked respondents the *Nation Hosting* measure alongside questions for a separate study. Of the 2,042 February respondents, 1,776 (87.0 percent) completed this survey. Section 1.3.1 compares these different samples. Enumerator identifiers are available for the February and May 2026 surveys but not the 2023 baseline. Section 5.1 shows the 2026 estimates are robust to enumerator fixed effects.

1.1.2 Panel

The face-to-face panel was administered by TIFA enumerators in Kakuma across three waves in 2025, with a fourth phone wave in February 2026 as part of the nationally representative survey, described in Section 3.4. Respondents were recruited through a multi-stage design. Sub-locations, the smallest administrative units, were treated as clusters and allocated sample using probability proportional to the number of households; within each cluster, households were selected at random; and within each household, a single respondent was drawn from the eligible members at random, with a further random selection where the first-drawn member was unavailable. Kenyan voters aged 18 to 65 resident in the study area were eligible. Where a panel respondent could not be re-interviewed, refresher respondents were added to maintain the cross-sectional sample. Enumerators were recruited from the Kakuma host community, which helped establish trust with participants, and interviewed the same respondents at each wave where possible. Section 5.1 verifies that the panel results are robust to enumerator fixed effects. Consent procedures, incentives, and data protections are described in Section 1.2.

1.1.3 Refugee-leader interviews

Alongside each face-to-face survey wave, I conducted semi-structured interviews with a panel of refugee leaders in Kakuma. Leaders were purposively sampled to vary by nationality, gender, and leadership role, including healthcare workers, local news broadcasters, leaders of refugee-led organizations, and employees of international organizations, and were recruited with the help of a local journalist and research assistant

with strong networks in Kakuma. Table 1 summarizes the composition of the interview panel, showing that it is broadly representative of the gender and nationality of the Kakuma/Kalobeyei refugee population.

Because UNHCR restrictions prevented in-person research in the camp during the cuts, interviews were conducted remotely over WhatsApp and lasted 30 to 60 minutes. Interviews were conducted in English. We first discussed the leader’s role in the camp. The topic guide then covered three areas. First, the cuts: which organizations and services were affected, the main impacts on refugees, coping strategies, and attribution of responsibility for the cuts. Second, integration and the Shirika Plan. Third, relations with the host community. The guide was held constant across waves so that changes in leaders’ accounts track changes in conditions, and each follow-up opened by asking what had changed since the previous interview. Consent procedures, compensation, and data protections are described in Section 1.2.

I reviewed the full set of interviews at each wave and selected quotations that illustrate themes recurring across multiple leaders, prioritizing accounts corroborated by other interviews or by the survey data. The interviews were not formally coded. Quotations are reproduced verbatim except for light edits for readability.

Table 1: Composition of refugee leader interview sample

Characteristic	Category	N	Sample (%)	Population (%)
Gender	Female	9	50.0	46.7
Gender	Male	9	50.0	53.3
Nationality	South Sudanese	10	55.6	61.7
Nationality	Somali	4	22.2	12.9
Nationality	Congolese	2	11.1	8.2
Nationality	Burundian	2	11.1	7.7

Note: Population percentages refer to the Kakuma/Kalobeyei refugee population (UNHCR operational data portal, March/April 2026).

1.2 Ethics

This research was conducted in accordance with APSA’s Principles and Guidance for Human Subjects Research. The study was approved by my university’s Institutional Review Board, and received Kenyan ethics approval from the Amref Health Africa Ethics and Scientific Review Committee (ESRC) and a research license from the National Commission for Science, Technology and Innovation (NACOSTI).

Phone-survey respondents gave verbal informed consent at the start of each call under a waiver of documentation, were told they could decline any question, and were asked to take the survey in a private place. Face-to-face panel respondents gave written consent before each interview, in English, Swahili, or Turkana according to their preference, and consent was re-obtained at each follow-up with a modified script stating that responses would be linked across waves. Refugee-leader interviewees gave verbal consent, could decline any question or end the interview at any time, and consented to being audio-recorded. Survey participants in the phone and face-to-face surveys received small phone airtime incentives, set by TIFA Research based on standard payments in the sector; refugee leaders received a payment of 1,000 KSh (approximately 8 USD) at each survey wave in mobile money, an amount advised by the research assistant on the basis of local norms. All survey data were anonymized by TIFA before being transferred to me. Refugee leaders are not identified in the paper, and their quotations were reviewed to remove potentially identifying details.

The nationwide surveys are representative of the Kenyan adult population and not composed mainly of vulnerable or marginalized groups. The Kakuma panel and refugee-leader interviews were conducted with individuals from relatively vulnerable populations, especially at the time of the cuts. Given this, the study was designed so that participation posed no more than minimal risk and questions about sensitive experiences could be declined. In particular, the interview questions with refugee leaders were designed to discuss the community as a whole, rather than their individual experiences. The research did not involve deception.

1.3 Sample

1.3.1 Nationwide surveys

Both nationwide surveys were conducted by phone, which is likely to skew toward respondents who are younger and more educated than the adult population. I therefore construct post-stratification weights for each survey using entropy balancing (Hainmueller 2012), which reweights each sample so that its distribution of age, gender, education, and region matches a target population distribution. Each survey is weighted to the Afrobarometer Kenya face-to-face survey conducted prior to that survey wave: the 2023 baseline to Round 9 (fielded 2021) and the 2026 follow-up to Round 10 (fielded 2024). I use Afrobarometer data rather than the census because the census does not include a measure of education. Weights are constructed separately for the 2023 and 2026 surveys, with the 2026 weights also applied to the May follow-up. All descriptive statistics and maps in the paper are weighted, and the difference-in-differences estimates are unweighted with HC1 standard errors.

I make three sample comparisons, reported in Table 2 descriptively and Tables 3 and 4 as formal tests. First, because the 2023 hosting items were asked only of the control arm, I compare that arm against the other arms to confirm the smaller control sample is not demographically distinct. The ‘Full nat-rep’ and ‘Control arm’ columns of Table 2 show the two are similar, and Panel A of Table 3 finds no covariate differs significantly between them.

Second, I compare the 2023 and 2026 nationally representative samples. The ‘2023 Full nat-rep’ and ‘2026 Full nat-rep’ columns of Table 2 show the two are similar, and Panel B confirms they match on gender and education; the one difference is age, with 2026 respondents on average 1.2 years older ($p < .01$). I account for this by including age, along with gender and education, as a control in all cross-survey specifications, and by weighting all descriptive statistics and maps. Third, I compare the May recontacted respondents against the full February sample. Of the 2,042 February respondents, 87.0 percent responded in May. The ‘Full nat-rep’ and ‘Follow-up’ columns of Table 2 show the recontact subsample is nearly identical to the February sample on every demographic. Table 4 regresses response to the May survey on demographics from the February survey wave. The only predictor of response is age, and its effect is small, with each additional year raising the response probability by 0.2 percentage points. As above, to account for this, I include age as a control in every specification and the weighting strategy. Importantly, there is no evidence that those who were more or less supportive of refugees were more likely to attrit from the study.

Table 2: Nationwide survey samples and external benchmark

	2023		2026		Afrobarometer
	Full nat-rep	Control arm	Full nat-rep	Follow-up	R10 (2024)
<i>N</i>	2510	816	2042	1776	2400
Age	36.92 (11.84)	36.67 (11.67)	36.07 (12.49)	36.42 (12.64)	36.88 (14.99)
Female	50.0%	47.2%	49.9%	49.9%	49.9%
Education: secondary or higher	48.3%	47.1%	68.7%	69.3%	55.8%
Region					
Central	10.3%	9.7%	11.7%	12.2%	12.6%
Coast	10.1%	10.1%	10.7%	10.5%	9.0%
Eastern	13.3%	13.3%	15.2%	13.8%	14.5%
Nairobi	9.8%	8.1%	9.0%	9.6%	10.2%
North Eastern	6.2%	7.4%	4.7%	4.7%	4.6%
Nyanza	12.2%	11.8%	12.4%	12.3%	12.7%
Rift Valley	25.5%	26.3%	26.4%	26.8%	26.0%
Western	12.6%	13.3%	9.9%	10.2%	10.4%

Note: Continuous variables: weighted mean (SD). Other rows: weighted percentages. Region percentages are weighted shares of each province and sum to 100 within each column.

Table 3: Tests of nationwide sample differences

Covariate	Diff	SE
Panel A: 2023 framing arms (other vs. refugees arm); joint F = 0.81, p = 0.488		
Age	-0.285	(0.492)
Female	-0.026	(0.021)
Education: secondary or higher	-0.010	(0.019)
Panel B: 2023 vs. 2026 nationally representative; joint F = 4.10, p = 0.006		
Age	1.164**	(0.361)
Female	0.007	(0.015)
Education: secondary or higher	0.001	(0.013)

Note: Each row reports the difference in means between the two samples with an HC1 standard error in parentheses. The joint test is an HC1-robust Wald test that all covariates are unrelated to sample membership. Tests are unweighted. Significance: *** p<.001, ** p<.01, * p<.05, † p<.10.

Table 4: Correlates of May 2026 follow-up survey retention

Predictor	Coef.	SE
Age	0.002***	(0.001)
Female	-0.000	(0.015)
Education: secondary or higher	0.029	(0.019)
Approval of national govt (Feb.)	-0.007	(0.005)
Integration index (Feb.)	0.013	(0.011)

Note: Linear probability model of responding to the May 2026 recontact on February 2026 observables, HC1 standard errors. Joint Wald test that all predictors are zero: F = 3.19, p = 0.007. N = 2,000. Significance: *** p<.001, ** p<.01, * p<.05, † p<.10.

1.3.2 Kakuma

Table 5: Kakuma sample: balanced panel vs. repeated cross-section

	Panel	Repeated cross-section		
		W1	W2	W3
<i>N</i>	205	300	301	301
Age	30.17 (9.14)	30.04 (8.88)	29.85 (8.92)	29.72 (8.84)
Female	48.8%	47.0%	47.2%	47.2%
Somali	0.5%	0.3%	0.0%	0.7%
Turkana (ethnicity)	95.1%	94.0%	95.0%	94.3%
Education: None	14.6%	14.3%	13.0%	14.6%
Education: Primary	17.6%	16.7%	18.6%	18.3%
Education: Secondary	38.5%	38.0%	40.5%	36.9%
Education: Tertiary	28.8%	30.7%	27.6%	25.9%
Monthly income (KES)	12,399 (26,120)	14,191 (21,299)	10,453 (20,681)	13,943 (35,079)
Food insecure	84.5%	86.9%	81.3%	79.7%
Approval of national govt (1–5)	2.98 (1.73)	2.84 (1.71)	3.00 (1.78)	3.13 (1.74)

Note: Continuous variables: mean (SD). Other rows: %. The Panel column is the balanced three-wave panel (respondents observed in all three waves), with time-invariant traits measured at Wave 1 and time-varying measures pooled across waves; W1–W3 are the full face-to-face waves, including refreshers added in each wave.

Table 6: Attrition selectivity on Wave 1 observables

	(1) LPM	(2) Logit
Constant	0.352** (0.124)	-0.575 (0.704)
Age	-0.003 (0.003)	-0.015 (0.017)
Female	-0.038 (0.054)	-0.192 (0.274)
Secondary or higher	0.013 (0.060)	0.071 (0.313)
W1 hosting support	0.015 (0.015)	0.075 (0.077)
Baseline beneficiary	-0.046 (0.055)	-0.228 (0.270)
N	300	300
Joint test	$F = 0.76$	$\chi^2 = 3.44$
Joint p	0.576	0.633

Note: The dependent variable indicates attriting after Wave 1. Column (1) is a linear probability model with HC1 robust SEs; column (2) a logit reporting raw coefficients with conventional SEs (in parentheses). The LPM joint test is a robust Wald F ; the logit joint test a likelihood-ratio χ^2 . Significance: *** $p < .001$, ** $p < .01$, * $p < .05$, † $p < .10$.

The Kakuma panel was conducted face-to-face, with the intention of surveying the same 300 host-community respondents at roughly two-month intervals. The Wave 1 cohort comprised 300 tracked respondents, of whom 217 (72%) completed the survey at Wave 2 and 205 at Wave 3; 83 attrited after Wave 1 and a further 12 between Waves 2 and 3. Waves 2 and 3 comprised 301 and 301 interviews respectively, the balance made up of refresher respondents recruited to maintain the repeated cross-section.

Attrition was concentrated between Waves 1 and 2, which coincided with a period of protests and heightened insecurity around Kakuma. Enumerators struggled to reach some households, and some respondents declined to be interviewed because of the unrest, producing a higher attrition rate than anticipated. Table 6 shows the results of regressing an indicator for attriting after Wave 1 on Wave 1 age, gender, education, hosting support, and baseline beneficiary status, the H2 treatment. None of the five predictors is individually significant, and the joint tests cannot reject that all five are simultaneously zero, so those who left the panel do not differ observably from those who stayed. However, note that these tests may be underpowered to detect changes.

Because no Wave 1 observable predicts three-wave retention (Table 6), inverse-probability-of-retention weighting should leave the intent-to-treat estimates essentially unchanged, and it does (Table 26). I estimate each retained respondent’s probability of appearing in all three waves from a logit on age, gender, education, hosting support, and beneficiary status, then weight the fixed-effects intent-to-treat regressions by the normalized inverse of the fitted probabilities. Adjusting for differences between respondents who remained in the panel and those who left does not change the main result. Note that this analysis corrects for selection on observables only. Attrition correlated with unobserved trajectories is addressed instead by the repeated cross-section replication (Section 3.3), which finds the same pattern on a sample refreshed at every wave.

1.4 Survey instruments

This section reports the English wording of every survey item used in the analysis. Response scales are named once here and referenced by shorthand after each item. Unless noted otherwise, every item carries a do-not-read “don’t know or refused” option, and in cleaning I reorient every outcome so that higher values denote more pro-refugee or pro-host responses.

- [SO] — five-point support scale: strongly support / somewhat support / neither support nor oppose / somewhat oppose / strongly oppose.
- [AD] — five-point agreement scale: strongly disagree (1) / somewhat disagree / neither agree nor disagree / somewhat agree / strongly agree (5).
- [YN] — yes / no.

1.4.1 Nationwide outcome items

Nation Hosting [SO]. “We now want to ask you more broadly about refugees. To what extent do you support or oppose Kenya hosting refugees?”

Local Hosting [SO]. 2023 Natrep: “Suppose the Kenyan government is going to enable refugees to move to different areas throughout the country. Would you support or oppose refugees moving into your constituency?” 2026 Natrep: “Please tell me how strongly you support or oppose the following policies: Refugees moving to live in your district.” The item was reworded slightly between rounds so that it read naturally alongside the 2026 integration battery, into which it was folded. The shift from “constituency” to “district” does not introduce a meaningful difference: “district” is a colloquial term for what became the electoral constituency after the 2010 constitutional reforms, and the terms are used interchangeably in Kenya. TIFA Research, the survey firm, confirmed the two are treated as equivalent. Also, because the difference-in-differences compares each round to itself across counties, any level shift common to all respondents from the rewording is absorbed by the Post main effect and does not enter the treatment interaction. Panel: Introduced with “Please tell me how strongly you support or oppose the following policies,” followed by the text: “Refugees continuing to live in your district”. The phrasing differs slightly to reflect the hosting-area context. Because the analysis compares changes within the same panel, this wording difference does not affect the estimates.

Policy preference (source of *Not Repat.*). “Which of these five statements is closest to your own opinion towards refugees living in Kenya?” Five options read in the following order: “(1) they should live in camps and not be able to legally work or move; (2) they should live in settlements and be able to legally work and move in Turkana and Garissa; (3) they should live throughout Kenya and be able to legally work and move freely; (4) they should be resettled to other countries; (5) they should be sent back to their country of origin.” *Not Repat.* is coded 1 for any response other than option (5) and 0 for option (5).

Refugee numbers (replicating Afrobarometer Round 10 Q73C verbatim). “Do you think Kenya should allow more or fewer refugees from other countries to come and live in this country?” [more refugees / about the same / fewer refugees / no refugees at all / don’t know or refused]. The indicator used in the interim-benchmark analysis (Section 2.7.2) equals 1 for “more refugees” or “about the same” and 0 otherwise.

Integration items [SO]. Introduced with “Please tell me how strongly you support or oppose the following policies,” followed by the text:

- Healthcare access: “Refugees having access to Kenyan doctors and hospitals and paying the same price as Kenyans.”
- Right to work: “Refugees being allowed to work in your community.”
- Freedom of movement: “Refugees being able to move freely and settle throughout Kenya.”

Camp closure (reversed as *Oppose Camp Closure*) [SO]. “Please tell me how strongly you support or oppose the following policies,” followed by: “The Kenyan government closing Kakuma camp and moving the refugees in Kakuma back to their countries.”

Hypothetical full-cut scenario [SO]. “Currently, Kenya receives financial support from international donors to host refugees. Imagine a situation where donors completely cut off/end their financial support for refugees in Kenya. This would mean that they no longer provide any of the aid that is used for services for refugees, such as healthcare and education. In this hypothetical situation, to what extent would you support Kenya continuing to host refugees?”

Empathy items [AD]. Index construction is described in Section 4.5.1. Introduced with “I’m going to now say a few statements. Please indicate how much you agree with each statement when you think about

refugees living in Kenya.” Presented in random order:

- “You are very concerned about the refugees and their suffering.”
- “You are worried that the refugees do not have enough food to eat because of the cuts to cash assistance and food rations.”

Trust in refugees [AD]. “To what extent do you agree or disagree with the following statement? ‘Generally speaking, most people in the refugee community are trustworthy.’”

Exposure items defining baseline-beneficiary status [YN].

- Camp healthcare use: “Have you or anyone in your household accessed hospitals, clinics, or dispensaries in Kakuma camp or Kalobeyi for any health issue in the past month, including the issue we asked about in the previous question or any other medical issue?”
- Camp education use: “Do you have any children in your household who have accessed schooling or education services in Kakuma camp or Kalobeyi in the past month?”
- NGO/aid employment: “Have you (or anyone in your household) been employed by an NGO in Kakuma camp or Kalobeyi in the past month, including casual work?”
- Employment by a refugee household: “Have you (or anyone in your household) been employed by a refugee (for things like housekeeping, cooking, childcare) in Kakuma camp or Kalobeyi in the past month?”
- Business serving refugees: “Do you or anyone in your household have a business that has sold goods or products to refugees in the past month?”

1.4.2 Open-ended questions and coding

Changes observed from the cuts (Kakuma, W1–W3). Respondents were first asked “Have you heard about recent cuts and reductions to services for refugees in Kakuma camp?” Those who answered yes were then asked, “What changes have you seen as a result of these cuts?”

Reasons for supporting or opposing camp closure (Kakuma, W3). Asked after the camp-closure question: “Why are you supportive/opposed/neutral toward the Kenyan government closing Kakuma camp and moving refugees back to their countries?” The wording was piped from the respondent’s previous answer.

Reasons for supporting or opposing hosting refugees in Kenya (Nationwide, 2023). Asked after the *Nation Hosting* question: “Why do you support/oppose Kenya hosting refugees?” The wording was piped from the respondent’s previous answer.

The unit of coding is the respondent-wave, and responses can receive multiple codes, so the category shares in Figure 5 of the paper sum to more than 100 percent. The three open-ended questions map to the figure as follows: the camp-closure reasons (Kakuma, W3) are shown in Panel A, for respondents who opposed closure ($N = 163$), and Panel B, for those who supported it ($N = 124$), split by the respondent’s answer to the closure item; the changes observed from the cuts (Kakuma, pooled W1–W3, $N = 615$ respondent-waves) shown in Panel C; and the reasons for hosting refugees (2023 national sample, control arm) shown in Panel D.

Coding proceeded in two stages. First, I developed the codebook inductively: for each question, I read the first 10 percent of responses and defined a category scheme, with labels matching those in Figure 5. Second, the remaining responses were coded against the fixed codebook by a large language model (Claude Opus 4.7, Anthropic). Before coding the full set, I provided the model with the category definitions and worked examples, then piloted the codebook on a subset of responses. I then revised the definitions where the model’s codes diverged from my own, iterating until the codebook produced reliable classifications. Responses of “don’t know,” refusals, and blank entries were excluded from the coded denominators.

1.4.3 Vignette experiment (W3)

Each respondent evaluated six scenarios drawn from the full $3 \times 2 \times 2 \times 2 = 24$ -profile design, presented in random order.

Introduction. “Next you will see several possible scenarios. For each scenario, please tell us how much you support or oppose Kenya hosting refugees. Use a scale from 1 (strongly oppose) to 5 (strongly support).”

Scenario template. “In this scenario, international organizations provide [AID] to refugees. Health care and schools are [ACCESS]. International organizations [JOBS]. In the countries refugees come from, [CONFLICT].”

The four attributes took the following level texts, each randomized independently:

Attribute	Levels (text shown to respondents)
Aid level	“100% of food rations and cash support (‘bamba chakula’)” — “40% of food rations and cash support (‘bamba chakula’)” — “0% (no food rations or cash support (‘bamba chakula’))”
Service access	“shared with the host community (refugees and hosts can both use them)” — “for refugees only”
Local jobs	“hire local residents” — “do not hire local residents”
Origin context	“conflict is ongoing” — “peace has been restored”

Example profile. “In this scenario, international organizations provide 40% of food rations and cash support (‘bamba chakula’). Health care and schools are for refugees only. International organizations hire local residents. In the countries refugees come from, conflict is ongoing.”

Outcome question (after each scenario). “On a scale from 1 (strongly oppose) to 5 (strongly support), how much do you support Kenya hosting refugees under this scenario?” [1 = strongly oppose / 2 = somewhat oppose / 3 = neither support nor oppose / 4 = somewhat support / 5 = strongly support].

Manipulation check (asked once, after the final scenario). “In the scenario you just saw, what was the level of food/cash support?” [100% / 40% / 0%].

1.5 Estimation equations

This section reports the estimating equations for analyses in the main paper and appendix. All models are estimated by ordinary least squares, with linear probability models for binary outcomes. Cross-sectional specifications use heteroskedasticity-robust (HC1) standard errors; panel specifications cluster standard errors by respondent.

1.5.1 Nationwide difference-in-differences (H1)

The pooled difference-in-differences in Table 2 of the paper is

$$Y_i = \alpha + \beta \text{Post}_i + \rho \text{Host}_i + \tau (\text{Post}_i \times \text{Host}_i) + \mathbf{X}_i' \Gamma + \varepsilon_i, \quad (1)$$

where Post_i indicates the 2026 round, Host_i indicates residence in a refugee-hosting county (Turkana or Garissa), and $\mathbf{X}_i$ contains age, gender, and an education index. The coefficient of interest is τ . Section 2.2 re-estimates equation (1) with the entropy-balancing post-stratification weights, Section 2.5 replaces HC1 inference with a wild cluster bootstrap clustered on county, and Section 2.6 varies the composition of the treated group. The Afrobarometer difference-in-differences in Section 2.7.2 estimates equation (1) with the 2024 Afrobarometer as the pre-period and with a different outcome variable.

1.5.2 Kakuma panel event study (H1)

The over-time analysis on the balanced three-wave panel (Table 3, Panel A of the paper) is

$$Y_{it} = \mu + \nu_i + \delta_2 D_t^2 + \delta_3 D_t^3 + \varepsilon_{it}, \quad (2)$$

with respondent fixed effects ν_i , wave indicators D_t^2 and D_t^3 (Wave 1 omitted), and standard errors clustered by respondent. The Wave 3 versus Wave 2 contrast is read from the same specification re-referenced to Wave 2. The repeated cross-section version (Section 3.3) replaces ν_i , which is unavailable once refresher respondents are included, with the demographic controls $\mathbf{X}_{it}$ from equation (1), retaining respondent-clustered standard errors:

$$Y_{it} = \mu + \delta_2 D_t^2 + \delta_3 D_t^3 + \mathbf{X}_{it}' \Gamma + \varepsilon_{it}. \quad (3)$$

The integration and trust outcomes (Section 3.5) and the four-wave extension (Section 3.4) use equation (2) with the corresponding outcomes and wave indicators.

1.5.3 Personal loss (H2)

All three tests of personal loss (discussed in Section 3.6.2) share the same estimator, a respondent-fixed-effects difference-in-differences on the balanced three-wave Kakuma panel, with standard errors clustered by respondent. They differ only in how they define which households are exposed to personal loss. Each takes the form

$$Y_{it} = \mu + \nu_i + \delta_2 D_t^2 + \delta_3 D_t^3 + \tau_2 (G_i \times D_t^2) + \tau_3 (G_i \times D_t^3) + \varepsilon_{it}, \quad (4)$$

with respondent fixed effects ν_i , wave indicators D_t^2 and D_t^3 (Wave 1 omitted), and an exposure indicator G_i interacted with each post wave. The interactions τ_2 and τ_3 give the differential change in exposed households' attitudes at the trough and at the partial-recovery wave. What changes across the three designs is the definition of G_i .

The baseline-beneficiary specification (Panel A) sets $G_i = B_i$, where B_i equals one if the household accessed camp healthcare or education or was employed by an NGO or a refugee in the 30 days before Wave 1. Because B_i is fixed before the cuts, it cannot be contaminated by attitudes that formed later, which makes this the intent-to-treat design and the primary test of H2. Table 3, Panel B of the paper reports the equivalent pooled parameterization, which replaces the respondent fixed effects with the beneficiary main effect ρB_i so that the pre-cut level gap between beneficiary and non-beneficiary households is displayed alongside the interactions:

$$Y_{it} = \mu + \delta_2 D_t^2 + \delta_3 D_t^3 + \rho B_i + \tau_2 (B_i \times D_t^2) + \tau_3 (B_i \times D_t^3) + \varepsilon_{it}. \quad (5)$$

In a fully balanced panel the fixed-effects and pooled parameterizations return identical interaction estimates; here they differ only through item-level nonresponse, and the estimates are nearly identical (Section 3.6.2).

The cohort difference-in-differences (Panel B) sets $G_i = L_i$, an indicator for first losing access to a benefit between Waves 2 and 3, estimated on the sample that excludes households that lost access between Waves 1 and 2. Here τ_2 is a placebo: because the loss occurs after Wave 2, the Wave 2 interaction should be near zero if exposed and unexposed households were on parallel trajectories beforehand, and τ_3 is the effect of the realized loss.

The switching difference-in-differences (Panel C) defines exposure by within-person loss events rather than a fixed group, pooling loss events across waves and domains:

$$Y_{it} = \nu_i + \lambda_t + \tau P_{it} + \varepsilon_{it}, \quad (6)$$

where P_{it} switches to one from the wave after household i first loses access in a domain, with respondent fixed effects ν_i , wave fixed effects λ_t , and respondent-clustered standard errors. The four domains (education, healthcare, refugee employment, and NGO employment) are estimated separately and pooled as loss in any domain, so τ recovers a single within-person effect of losing access.

1.5.4 Vignette experiment (H1)

Each respondent i rates six profiles j ; the average marginal component effects in Figure 4 of the paper come from a single linear model of rated support on all attribute-level indicators,

$$Y_{ij} = \nu_i + \sum_a \sum_{\ell \neq \ell_0(a)} \beta_{a\ell} T_{ij}^{a\ell} + \varepsilon_{ij}, \quad (7)$$

where $T_{ij}^{a\ell}$ indicates that profile j presents level ℓ of attribute a (reference levels ℓ_0 : aid at 100 percent, services shared, hire locals, conflict ongoing), with respondent fixed effects and standard errors clustered by respondent. The empathy moderation (Table 31, discussed in Section 4.5) interacts every attribute with the standardized two-item empathy index E_i :

$$Y_{ij} = \nu_i + \sum_a \sum_{\ell \neq \ell_0(a)} [\beta_{a\ell} T_{ij}^{a\ell} + \pi_{a\ell} (T_{ij}^{a\ell} \times E_i)] + \varepsilon_{ij}, \quad (8)$$

so that $\beta_{a\ell}$ is the contrast at mean empathy and $\pi_{a\ell}$ its change per standard deviation of empathy. Because E_i is absorbed by the respondent fixed effects, its main effect is not separately estimated. Marginal means by empathy group appear in Section 4.5.2.

1.6 Hosting Index construction

I combine the hosting items into a single index rather than analyzing each separately to avoid the multiple-comparisons problem that arises from testing several correlated outcomes one at a time. A summary index controls the family-wise error rate and, by pooling the shared signal across items, typically yields a more powerful test than any single item (Anderson 2008). I use Anderson’s inverse-covariance weighting, which down-weights near-redundant components so that highly correlated items do not dominate the index.

The three items measure related but distinct facets of support for hosting: support for hosting nationally, support for hosting locally, and rejection of repatriation. Combining them requires only that they share a common orientation, which holds: all three pairwise correlations are positive in both the Kakuma panel and the nationwide baseline, with mean absolute correlations of 0.43 and 0.51 (Table 7), and Cronbach’s alpha is 0.69 and 0.67. The weighting choice makes little difference here: because the components are similarly and only moderately correlated, the implied ICW weights are nearly equal, and the index correlates with a simple standardized average at 1.00 in the panel and 0.997 nationwide. The main results are unchanged under the simple average (Section 2.7.1).

Table 7: Hosting Index diagnostics

Statistic	Kakuma panel	Nationwide
Mean r among items	0.43	0.51
Cronbach’s alpha	0.69	0.67
ICW vs. simple-average correlation	1.000	0.997

Note: Mean absolute pairwise correlation among the three index components, Cronbach’s alpha, and the correlation between the inverse-covariance-weighted index and a simple standardized average. Panel components: Nation Hosting, Local Hosting, Oppose Closure. Nationwide components: Nation Hosting, Local Hosting, Not Repat.

2 Nationwide results and robustness

2.1 Summary statistics

Table 8 reports the distribution of the four outcomes by round for the national sample and the Turkana subsample. Nationally, all four outcomes increase or hold steady between 2023 and 2026: the hosting index rises from -0.12 to 0.05 , and the share not selecting repatriation rises from 64 to 76 percent. In Turkana, by contrast, the index falls from 0.24 to 0.09 and local hosting support falls from 3.31 to 2.98 , while national hosting support and the repatriation measure hold roughly constant. The table also shows that Turkana was more pro-hosting than the national average on every measure in 2023, a gap that closes by 2026.

Table 8: Nationwide refugee attitudes: outcome distributions by round

Outcome	Year	N	Mean	SD	P25	Med.	P75	Min	Max
National									
Hosting Index	2023	814	-0.12	1.05	-0.80	0.13	0.86	-1.83	1.23
	2026	2,041	0.06	0.86	-0.38	0.08	0.81	-2.04	1.37
Nation Hosting	2023	809	3.55	1.59	2.00	4.00	5.00	1.00	5.00
	2026	1,776	3.62	1.61	2.00	4.00	5.00	1.00	5.00
Local Hosting	2023	810	2.99	1.76	1.00	4.00	5.00	1.00	5.00
	2026	2,028	3.11	1.76	1.00	4.00	5.00	1.00	5.00
Not Repat.	2023	801	0.64	0.48	0.00	1.00	1.00	0.00	1.00
	2026	2,018	0.77	0.42	1.00	1.00	1.00	0.00	1.00
Turkana									
Hosting Index	2023	150	0.24	0.91	-0.18	0.60	1.07	-1.83	1.07
	2026	349	0.09	0.94	-0.33	0.03	1.07	-1.90	1.37
Nation Hosting	2023	150	4.15	1.33	4.00	5.00	5.00	1.00	5.00
	2026	72	3.94	1.61	4.00	5.00	5.00	1.00	5.00
Local Hosting	2023	150	3.31	1.78	1.00	4.00	5.00	1.00	5.00
	2026	346	2.98	1.78	1.00	3.50	5.00	1.00	5.00
Not Repat.	2023	148	0.80	0.40	1.00	1.00	1.00	0.00	1.00
	2026	343	0.80	0.40	1.00	1.00	1.00	0.00	1.00

Note:

National means and standard deviations are weighted (entropy-balancing post-stratification weights); Turkana means and standard deviations are unweighted, since the post-stratification weights target the national population. Percentiles, minimum, and maximum are unweighted throughout.

2.2 Weighted difference-in-differences analysis

The difference-in-differences estimates in the main text are unweighted with HC1 standard errors, as pre-specified (PAP Section 4.4). Table 9 re-estimates all four specifications applying the entropy-balancing post-stratification weights described in Section 1.3. The weights are constructed for the nationally representative samples, so the 2026 Turkana oversample carries no post-stratification weight. Dropping these respondents would reduce the treated post-period cell from 383 to 118 and discard most of the treated-group information in the weighted specification. I therefore retain the oversample, assigning its respondents unit weight and normalizing weights to mean one within each period-by-hosting-county cell so that the two weight systems sit on a common scale. This keeps all treated respondents in the weighted estimation while weighting the nationally representative respondents by their post-stratification weights. Point estimates are stable across panels, and the *Hosting Index*, *Local Hosting*, and *Not Repat.* effects remain significant; only *Nation Hosting* loses significance under the weights, reflecting the larger standard errors that weighting introduces rather than a change in sign or magnitude.

2.3 Full estimation results for the main difference-in-differences

Table 10 reports the full estimation results for the pooled difference-in-differences presented in Table 2 of the main paper. Each column is a separate unweighted OLS model of the named dependent variable, with HC1 standard errors, estimated on the identical sample and specification as the main text; the treatment coefficients reproduce Table 2 by construction. The control coefficients follow familiar patterns from research on immigration attitudes: older respondents are less supportive of refugees on every outcome, more educated respondents are more supportive of hosting nationally and less likely to prefer repatriation, and gender is unrelated to any outcome. Notably, education does not predict support for hosting refugees locally, consistent with the paper’s broader distinction between support for the principle of hosting and willingness to host nearby.

Table 9: Nationwide difference-in-differences, unweighted and weighted

V1	Hosting Index	Nation Hosting	Local Hosting	Not Repat.
Panel A: Unweighted (main specification)				
Post	0.114** (0.040)	-0.027 (0.066)	0.099 (0.072)	0.079*** (0.019)
Hosting county	0.326*** (0.063)	0.562*** (0.095)	0.381** (0.120)	0.109*** (0.028)
Post x Hosting county	-0.365*** (0.081)	-0.418* (0.196)	-0.659*** (0.155)	-0.099** (0.036)
Observations	3,401	2,841	3,381	3,357
Panel B: Entropy-balancing weights				
Post	0.166*** (0.050)	0.042 (0.081)	0.160† (0.086)	0.101*** (0.023)
Hosting county	0.397*** (0.071)	0.674*** (0.109)	0.460*** (0.133)	0.135*** (0.032)
Post x Hosting county	-0.413*** (0.088)	-0.390† (0.204)	-0.708*** (0.168)	-0.120** (0.040)
Observations	3,401	2,841	3,381	3,357

Note:

OLS with HC1 standard errors in parentheses; demographic controls (age, gender, education) included in all models. Panel B weights observations by the entropy-balancing post-stratification weight for each survey round. All outcomes coded so higher values are more pro-refugee. Significance: *** p<.001, ** p<.01, * p<.05, † p<.10.

Table 10: Nationwide difference-in-differences: full estimation results

V1	Hosting Index	Nation Hosting	Local Hosting	Not Repat.
Post	0.114** (0.040)	-0.027 (0.066)	0.099 (0.072)	0.079*** (0.019)
Hosting county	0.326*** (0.063)	0.562*** (0.095)	0.381** (0.120)	0.109*** (0.028)
Post × Hosting county	-0.365*** (0.081)	-0.418* (0.196)	-0.659*** (0.155)	-0.099** (0.036)
Age	-0.011*** (0.001)	-0.006* (0.003)	-0.025*** (0.003)	-0.003*** (0.001)
Gender	-0.003 (0.031)	0.017 (0.058)	0.082 (0.060)	-0.019 (0.015)
Education index	0.083*** (0.020)	0.146*** (0.037)	-0.016 (0.037)	0.061*** (0.009)
Constant	0.098 (0.100)	3.427*** (0.183)	3.944*** (0.187)	0.616*** (0.047)
Observations	3,401	2,841	3,381	3,357
R ²	0.040	0.023	0.034	0.033

Note: Each column is a separate unweighted OLS model of the named dependent variable, HC1 standard errors in parentheses. Outcomes coded so higher values are more pro-refugee. Significance: *** p<.001, ** p<.01, * p<.05, † p<.10.

2.4 Covariate-by-period interactions

The additive specification in equation 1 assumes that the association between each demographic control and the outcome is constant across survey rounds. Because the 2023 and 2026 samples differ slightly in age (Section 1.3), I probe this assumption by interacting each demographic control, centered at the estimation-sample mean, with the Post indicator. This allows the demographic gradients to differ by round while leaving the Post × Hosting county coefficient interpretable as the differential change in hosting counties at

mean covariates. Table 11 reports the estimates alongside the additive baseline. The treatment estimates are stable: the Post $\times$ Hosting county coefficients are slightly larger in magnitude than in the additive specification on all four outcomes and retain the same significance levels, so the main results are, if anything, conservative with respect to compositional change in the demographic gradients between rounds.

Table 11: Nationwide difference-in-differences with covariate-by-period interactions

	Hosting Index	Nation Hosting	Local Hosting	Not Repat.
Panel A: Additive controls (main specification)				
Post	0.114** (0.040)	-0.027 (0.066)	0.099 (0.072)	0.079*** (0.019)
Hosting county	0.326*** (0.063)	0.562*** (0.095)	0.381** (0.120)	0.109*** (0.028)
Post $\times$ Hosting county	-0.365*** (0.081)	-0.418* (0.196)	-0.659*** (0.155)	-0.099** (0.036)
Observations	3,401	2,841	3,381	3,357
R ²	0.040	0.023	0.034	0.033
Panel B: Demographics interacted with Post				
Post	0.120** (0.040)	-0.020 (0.066)	0.109 (0.072)	0.082*** (0.019)
Hosting county	0.355*** (0.064)	0.585*** (0.096)	0.430*** (0.120)	0.118*** (0.028)
Post $\times$ Hosting county	-0.398*** (0.082)	-0.448* (0.196)	-0.721*** (0.157)	-0.108** (0.036)
Age	-0.012*** (0.003)	-0.009* (0.004)	-0.023*** (0.005)	-0.004** (0.001)
Female	0.064 (0.059)	0.072 (0.090)	0.163 (0.105)	0.010 (0.027)
Education index	0.159*** (0.036)	0.208*** (0.054)	0.120† (0.062)	0.083*** (0.016)
Post $\times$ Age	0.001 (0.003)	0.003 (0.005)	-0.004 (0.006)	0.002 (0.001)
Post $\times$ Female	-0.100 (0.069)	-0.093 (0.118)	-0.122 (0.128)	-0.044 (0.032)
Post $\times$ Education index	-0.120** (0.043)	-0.112 (0.073)	-0.215** (0.078)	-0.036† (0.020)
Constant	-0.026 (0.035)	3.657*** (0.053)	3.047*** (0.061)	0.703*** (0.016)
Observations	3,401	2,841	3,381	3,357
R ²	0.044	0.025	0.036	0.036

Note: OLS with HC1 standard errors in parentheses; each column is a separate model of the named dependent variable. Panel A is the additive specification; Panel B centers each demographic control at the estimation-sample mean and interacts it with Post. Outcomes coded so higher values are more pro-refugee. Significance: *** $p < .001$, ** $p < .01$, * $p < .05$, † $p < .10$.

2.5 Wild cluster bootstrap

Because treatment is assigned at the county level and only two counties are treated (Turkana and Garissa), the analytic cluster-robust standard errors risk overstating precision. Table 12 therefore recomputes inference for the interaction term using a wild cluster bootstrap with Rademacher weights, clustering on county and referencing a t distribution with $G - 1$ cluster degrees of freedom. The *Hosting Index*, *Local Hosting*, and *Not Repat.* effects all remain significant under this more demanding standard, with bootstrap confidence intervals that exclude zero. Only the *Nation Hosting* estimate attenuates, from conventional significance to marginal ($p = 0.078$). This is consistent with the panel’s finding that there is most backlash to refugee hosting locally.

Table 12: Wild cluster bootstrap inference for the nationwide difference-in-differences

	Hosting Index	Nation Hosting	Local Hosting	Not Repat.
Post $\times$ Hosting county	-0.365	-0.418	-0.659	-0.099
Analytic (HC1) p	0.000	0.033	0.000	0.006
Bootstrap SE	(0.092)	(0.232)	(0.209)	(0.020)
Bootstrap p	0.000***	0.078†	0.003**	0.000***
Bootstrap 95% CI	[-0.551, -0.180]	[-0.885, 0.049]	[-1.080, -0.239]	[-0.140, -0.057]
Clusters	49	49	49	49

Note: Wild cluster bootstrap standard errors (Rademacher weights, 9,999 replications) clustered on county. Bootstrap p-values and confidence intervals reference a t distribution with G-1 cluster degrees of freedom. Analytic p-values are two-sided from HC1 standard errors. Significance from the bootstrap p-value: *** < .001, ** < .01, * < .05, † < .10.

2.6 County leave-one-out analysis

The treated group in the nationwide difference-in-differences pools two counties whose exposure and measurement differs: Turkana hosts Kakuma and is measured through a 2026 oversample, while Garissa hosts Dadaab, has a majority-Somali host population co-ethnic with most Dadaab refugees, and is included in the 2026 round through the nationally representative sample only. Table 13 therefore re-estimates equation (1) varying the composition of the treated group: excluding Garissa (Turkana as the sole treated county) and excluding Turkana (Garissa as the sole treated county), dropping the excluded county’s respondents from the estimation sample entirely. As a complementary check on the control group, I also re-estimate the *Hosting Index* specification dropping each control county one at a time and report the range of the resulting interaction estimates in the table note.

The *Hosting Index* and *Local Hosting* interactions remain negative and significant whether Turkana or Garissa is the sole treated county, so neither drives the pooled result. The Turkana-only estimates are close to the pooled ones, as expected given Turkana’s larger share of the treated sample; the Garissa-only estimates rest on far fewer treated post-period observations, since Garissa received no 2026 oversample, and their magnitudes should be read with that lower precision in mind. Dropping each control county in turn leaves the *Hosting Index* interaction in a tight band around the pooled estimate, so no single control county drives the comparison group either.

Table 13: Nationwide difference-in-differences: leave-one-out by treated county

Post $\times$ Hosting county	Both (main)	Excl. Garissa	Excl. Turkana
Hosting Index	-0.365*** (0.081)	-0.252** (0.096)	-0.640*** (0.171)
Nation Hosting	-0.418* (0.196)	-0.157 (0.228)	-0.925** (0.353)
Local Hosting	-0.659*** (0.155)	-0.424* (0.189)	-1.388*** (0.338)
Not Repat.	-0.099** (0.036)	-0.071† (0.043)	-0.126 (0.081)
N (treated, 2026)	383	349	34
Observations	3,401	3,216	2,902

Note: Each cell reports the Post $\times$ Hosting county interaction from the nationwide difference-in-differences, OLS with HC1 standard errors in parentheses and demographic controls (age, gender, education). Columns vary the treated group: both hosting counties, Turkana only, and Garissa only. Sample sizes are for the hosting-index specification. Dropping each of the 47 control clusters one at a time leaves the hosting-index interaction between -0.375 and -0.358 (pooled estimate -0.365; the interaction is estimable for 46 of the 47 drops). Significance: *** $p < .001$, ** $p < .01$, * $p < .05$, † $p < .10$.

2.7 Additional dependent variables

2.7.1 Hosting index robustness

The *Hosting Index* used in the main text is an inverse-covariance-weighted (ICW) composite of three five-point and binary items that mean-imputes any missing component, so respondents who answered only some of the three items still contribute. I check that the nationwide result is not an artifact of these construction choices. Table 14 re-estimates the pooled difference-in-differences with three alternative dependent variables alongside the main index: a complete-case ICW index, restricted to respondents with all three components present; a simple average of the three standardized components in place of the ICW weighting; and an ICW index built from binarized components, recoding each five-point item as an indicator for support before constructing the index.

All three alternatives return the same conclusion as the main table: the $\text{Post} \times \text{Hosting county}$ coefficient stays negative, significant at the same level, and comparable in magnitude. If anything, the complete-case restriction yields a slightly steeper decline, so the result is not driven by respondents with partial responses.

Table 14: Nationwide difference-in-differences with alternative index constructions

	Hosting Index (ICW)	Hosting Index (ICW, complete cases)	Hosting Index (simple average)	Hosting Index (binary components)
Post	0.114** (0.040)	0.096* (0.041)	0.101* (0.040)	0.107** (0.040)
Hosting county	0.326*** (0.063)	0.331*** (0.063)	0.337*** (0.063)	0.301*** (0.062)
Post $\times$ Hosting county	-0.365*** (0.081)	-0.445*** (0.107)	-0.363*** (0.081)	-0.333*** (0.079)
Controls	Yes	Yes	Yes	Yes
Observations	3,401	2,795	3,401	3,401
R ²	0.040	0.046	0.039	0.039

2.7.2 Afrobarometer refugee-numbers measure

The nationwide difference-in-differences spans a long window, from 2023 to 2026, over which attitudes could have shifted for reasons unrelated to the cuts. As an interim benchmark, I use a refugee-numbers item from Afrobarometer Round 10, fielded in mid-2024, after the initial WFP ration reductions but before the major 2025 cuts, and replicated verbatim in the 2026 follow-up. The item asks whether Kenya should host more refugees, about the same number, fewer, or none; I code an indicator for preferring the same number or more. Table 15 reports response shares by round, and Table 16 estimates the pooled difference-in-differences of equation 1, with the 2024 Afrobarometer sample as the pre period. The specification includes age, gender, and education, harmonized to the survey coding as in the weighting procedure (Section 1.3).

Two findings emerge. First, the national trend matches the main analysis: between 2024 and 2026, the share of Kenyans preferring that the country host the same number of refugees or more nearly doubled, from 19.9 to 38.0 percent ($\hat{\beta} = 0.177$, $p < 0.001$). The rise in support outside the hosting regions documented in Table 2 of the main paper therefore appears on an independently fielded measure over a shorter, post-2024 window, and is not an artifact of the 2023 baseline. Second, hosting counties moved in parallel with the rest of the country on this measure ($\hat{\tau} = 0.004$, s.e. 0.067). This pattern is consistent with the integration results below and with the main finding: exposed counties track the rest of the country on support for hosting refugees in Kenya broadly, while the decline documented in the main text is specific to willingness to host refugees locally. However, this null should be read with caution: Afrobarometer includes only 78 hosting-county respondents, so the minimum detectable effect for the interaction is roughly 18 percentage points.

Table 15: Preferred number of refugees hosted, 2024–2026

	National		Hosting counties	
	2024	2026	2024	2026
None	20.2	14.0	17.4	8.0
Fewer	59.9	48.5	66.4	57.4
About the same	9.6	18.9	4.8	16.6
More	10.3	18.6	11.3	18.1
Same or more (%)	19.9	37.5	16.1	34.7
N	2,369	1,757	78	100

Note: Weighted response shares. 2024: Afrobarometer Round 10; 2026: follow-up survey. Hosting counties are Turkana and Garissa.

Table 16: Difference-in-differences: preferred refugee numbers, 2024–2026

	Same/More refugees
Post	0.177*** (0.015)
Hosting county	-0.033 (0.046)
Post $\times$ Hosting county	0.004 (0.067)
Age	-0.002** (0.000)
Female	-0.005 (0.014)
Education index	-0.015† (0.009)
Constant	0.305*** (0.039)
N (hosting counties, 2024)	78
N (hosting counties, 2026)	100
Observations	4,112
R ²	0.040

Note: OLS with HC1 standard errors. Outcome is an indicator for preferring that Kenya host the same number of refugees or more. Pre period is the 2024 Afrobarometer; post is the 2026 follow-up. Controls: age, gender, education. *** $p < 0.001$, ** $p < 0.01$, * $p < 0.05$, † $p < 0.10$.

2.7.3 Integration measures

The 2026 survey includes measures of support for refugee integration: refugees’ access to healthcare, right to work, and freedom of movement, each on a five-point scale. Table 17 reports weighted national means and compares Turkana with the rest of the country. Nationwide support is moderately high, greatest for healthcare, then right to work, then freedom of movement. Turkana respondents are statistically indistinguishable from the rest of the country on all three items; the point estimates are small and negative and none approaches significance. That exposed and non-exposed counties now look alike is consistent with the convergence documented in the main text. MacDonald and Lichtenheld (2025) find that support for integration was higher in refugee-hosting counties in 2023, but by 2026 that difference between refugee-hosting counties and the rest of the country is no longer evident. I read this as suggestive rather than as a within-design estimate, since it compares a 2026 cross-section against a separately fielded 2023 study rather than tracking

the same integration items over time. The broader pattern is the same across measures: exposure to refugees in these counties no longer predicts higher support than the rest of the country.

Table 17: Support for refugee integration, 2026 follow-up

Outcome	National (wt.)	Rest of Kenya	Turkana	Difference	N
Access to healthcare	4.06	4.08	4.02	-0.07 (0.09)	2,315
Right to work	3.48	3.51	3.36	-0.15 (0.10)	2,320
Freedom of movement	3.12	3.13	3.00	-0.13 (0.11)	2,322

Note: Five-point items, higher values more supportive. National column: weighted mean on the nationally representative sample. Remaining columns: OLS of each outcome on a Turkana indicator, unweighted on all 2026 respondents including the oversample, HC1 standard errors in parentheses. *** $p < 0.001$, ** $p < 0.01$, * $p < 0.05$, † $p < 0.10$.

2.8 Alternative estimators

Beyond the index, I probe the two five-point outcomes individually with alternative estimators. Table 18 reports three specifications for each item: Panel A repeats the OLS estimate from the main table; Panel B recodes the outcome as a binary indicator for support, a rating of four or five, and estimates a linear probability model; and Panel C estimates an ordered logit on the full five-point scale, relaxing the assumption that the response categories are equally spaced.

The *Local Hosting* result is robust across all three. For *Nation Hosting*, the ordered logit matches the OLS result, but the binary recoding is no longer significant, indicating that the *Nation Hosting* movement occurs partly within the oppose side of the scale rather than across the support threshold. This is consistent with the wild cluster bootstrap results above: the *Nation Hosting* estimate is the least robust of the four outcomes.

Table 18: Nationwide difference-in-differences: binary recoding and ordered logit

	Nation Hosting	Local Hosting
<i>Panel A: OLS (five-point outcome)</i>		
Post $\times$ Hosting county	-0.418* (0.196)	-0.659*** (0.155)
<i>Panel B: LPM, binary outcome (support ≥ 4)</i>		
Post $\times$ Hosting county	-0.088 (0.055)	-0.174*** (0.044)
<i>Panel C: Ordered logit (latent scale)</i>		
Post $\times$ Hosting county	-0.587* (0.243)	-0.776*** (0.166)
Observations	2,841	3,381

3 Kakuma panel results and robustness

3.1 Summary statistics

Table 19 reports descriptive statistics for the three-wave Kakuma panel, wave by wave and pooled. It covers the four hosting outcomes and the ten household outcomes plotted in Figure 3 of the main text: eight service-access, employment, and security items, together with two economic measures, the share of households earning under KSh 5,000 a month and the share whose household head is unemployed. Two further household measures are reported here but not plotted: food insecurity of any severity, and median monthly earnings. Continuous and index measures are reported with their pooled standard deviation, minimum, and maximum.

Among the attitude measures, opposition to camp closure falls from 3.39 at Wave 1 to 3.13 at the Wave 2 trough and does not recover by Wave 3 (3.09). Local hosting support is flat. Support for hosting refugees

in Kenya nationally, by contrast, dips at Wave 2 (3.12 to 2.96) and then rises above its baseline by Wave 3 (3.40), and the hosting index correspondingly returns to its Wave 1 level (0.00, -0.10 , -0.02). The household measures show the material side of the same period. The Figure 3 outcomes trace the sustained decline in camp-area services and employment alongside the rise in feeling unsafe at night. On the economic measures, median household earnings fall from KSh 9,000 at Wave 1 to KSh 5,000 at Waves 2 and 3, and the share of households earning under KSh 5,000 a month rises from 33.9 to 44.4 and then 47.3 percent; neither recovers by Wave 3. Household-head unemployment peaks at Wave 2 (50.7 percent, against 42.0 at Wave 1 and 43.9 at Wave 3), and food insecurity remains high throughout.

Table 19: Kakuma panel descriptive statistics by wave

	By wave				Pooled		
	W1	W2	W3	All	SD	Min	Max
Attitudes toward refugees							
Nation Hosting (1–5)	3.12	2.96	3.40	3.16	1.76	1.00	5.00
Local Hosting (1–5)	3.11	3.10	3.03	3.08	1.78	1.00	5.00
Oppose camp closure (1–5)	3.39	3.13	3.09	3.20	1.82	1.00	5.00
Hosting Index (standardized)	0.00	-0.10	-0.02	-0.04	0.99	-1.55	1.26
Household outcomes							
Food insecure (any)	87.7	83.7	81.9	84.5			
Feels unsafe at night	58.1	69.3	71.6	66.3			
Household earnings under KSh 5,000	33.9	44.4	47.3	42.3			
Household head unemployed	42.0	50.7	43.9	45.5			
Crime victim in past month	52.7	52.2	42.9	49.3			
Uses camp healthcare	42.4	30.2	26.8	33.2			
Uses camp education	20.0	20.0	17.1	19.0			
Household has NGO job	19.5	17.1	14.1	16.9			
Household has refugee job	13.2	11.7	14.1	13.0			
Severe food insecurity	12.3	17.2	13.2	14.2			
Business serving refugees	13.2	16.6	13.2	14.3			
Median monthly income (KES)	9,000	5,000	5,000	5,000	26,120	0	300,000

Note: Balanced three-wave Kakuma panel: 205 respondents per wave (615 respondent-waves), unweighted. Row-level N ranges from 581 to 615 with item non-response. Binary rows are percentages. Attitude items are five-point, higher values more pro-refugee; Oppose camp closure reverses the camp-closure item. The Hosting Index is a standardized inverse-covariance-weighted (Anderson 2008) combination of the three attitude items, fit on Wave 1. Severe food insecurity is missing meals on 5–7 days in the past week; feels unsafe at night is disagreeing with feeling safe; service, employment, business and crime items use a one-month recall. Income is total household earnings, reported as a median. At Wave 1 a refusal option was recorded and is treated as missing (34 of 205), so the Wave 1 income figures rest on a smaller sample and are likely upward-biased. Shares are ordered by Wave 3 prevalence; wave-to-wave tests appear in Section 3.2.

3.2 Over-time welfare tests

Figure 3 in the main text plots the share of Kakuma panel households reporting each livelihood, security, service-access, and economic outcome at each wave. This section reports the formal tests behind that figure. For each of the ten outcomes I estimate a linear probability model of the binary outcome on wave indicators, with standard errors clustered by respondent, and record the wave-to-wave differences (W2–W1, W3–W1, W3–W2), a joint test that the three waves are equal, and the implied prevalence at each wave. Consistent with the small per-wave sample (roughly 200 balanced-panel households), most movements are imprecise. Rows are ordered by Wave 3 prevalence to match the figure. The two economic rows are the share of households earning under KSh 5,000 a month and the share whose household head is unemployed; the earnings row rests on a smaller Wave 1 sample because the Wave 1 instrument recorded a refusal option that is treated as missing, as described in the note to Table 19.

Table 20: Host livelihood outcomes across the Kakuma panel

Outcome	Prevalence (%)			Wave difference (pp)			Joint p
	W1	W2	W3	W2–W1	W3–W1	W3–W2	
Feels unsafe at night	58.1	69.3	71.6	+11.1** (4.1)	+13.4** (4.4)	+2.3 (4.3)	0.003
Household earnings under KSh 5,000	33.9	44.4	47.3	+10.5* (4.3)	+13.4** (4.3)	+2.9 (4.2)	0.005
Household head unemployed	42.0	50.7	43.9	+8.8* (4.3)	+2.0 (4.6)	-6.8 (4.7)	0.106
Crime victim in past month	52.7	52.2	42.9	-0.5 (3.9)	-9.8* (4.3)	-9.3* (4.1)	0.037
Uses camp healthcare	42.4	30.2	26.8	-12.2** (3.9)	-15.6*** (4.5)	-3.4 (4.1)	0.001
Uses camp education	20.0	20.0	17.1	+0.0 (3.1)	-2.9 (3.5)	-2.9 (3.3)	0.623
Household has refugee job	13.2	11.7	14.1	-1.5 (2.4)	+1.0 (3.0)	+2.4 (3.0)	0.683
Household has NGO job	19.5	17.1	14.1	-2.4 (3.1)	-5.4 (3.5)	-2.9 (3.2)	0.305
Severe food insecurity	12.3	17.2	13.2	+5.0 (3.2)	+1.0 (3.2)	-4.0 (3.1)	0.243
Business serving refugees	13.2	16.6	13.2	+3.4 (2.9)	+0.0 (3.1)	-3.4 (3.3)	0.444

Note: Each row is a linear probability model of a binary household outcome (coded 0/100) on wave indicators, balanced Kakuma panel, standard errors clustered by respondent. Differences are in percentage points. Joint p is a Wald test that the three waves are equal. Rows ordered by Wave 3 prevalence. Significance: *** $p < .001$, ** $p < .01$, * $p < .05$, † $p < .10$.

3.3 Repeated cross-section

The panel results section in the main paper analyzes the 205 respondents interviewed in all three waves, which isolates within-person change but discards the refresher respondents recruited at Waves 2 and 3 to replace attriters. Here I reproduce the over-time analysis on the full repeated cross-section. This trades the panel’s within-person identification for a larger sample, and serves as a check that the panel results are not an artifact of who remained in the study.

Because the repeated cross-section changes composition across waves — refreshers recruited at Waves 2 and 3 replace attriters, and attrition, while unrelated to Wave 1 observables (Section 1.3.2), was concentrated in the W1–W2 window — raw wave means could conflate genuine attitude change with compositional drift. I therefore estimate covariate-adjusted wave contrasts, regressing each outcome on wave indicators and controls for age, gender, and education (equation (3)), with standard errors clustered on respondent to account for the repeat interviews of panel members who reappear across cross-sections. In place of the balanced panel’s respondent fixed effects, the covariate adjustment holds the demographic composition of each wave fixed. Table 21 reports the results, using the same three outcomes and the same two contrasts (Wave 2 vs. Wave 1, and Wave 2 vs. Wave 3) as the balanced-panel event study in the main text (Table 3, Panel A of the paper).

The results mirror the balanced panel. Opposition to camp closure falls from Wave 1 to Wave 2 by 0.332 points ($p < 0.05$), a steeper and now conventionally significant version of the marginal decline in the panel, and does not recover by Wave 3. Local hosting support shows no significant over-time change. *Nation Hosting* support decreases significantly at Wave 2 (-0.275, $p < .05$) before rising over the stabilization period (0.629, $p < .001$ from Wave 2 to Wave 3) to end above its baseline. The main results therefore hold when including refresher respondents in the sample.

3.4 Fourth wave: longer-term changes

This subsection extends the panel to a fourth wave, fielded by phone in February 2026 alongside the nationwide follow-up survey, roughly five months after Wave 3. The design captures the period of partial aid reinstatement described in Section 3.2 of the main paper: between Waves 3 and 4, rations were raised for Categories 1 through 3 and cash assistance resumed, though not to pre-cut levels. Wave 4 therefore extends the observation window from the acute-crisis phase into a period of partial recovery, allowing me to ask whether the Wave 2–3 decrease in support persisted as conditions stabilized.

Two features of Wave 4 warrant caution, and I read its results as suggestive of longer-run trajectories rather than as confirmatory. First, the switch from face-to-face to phone administration introduces a change in survey mode that could shift responses independently of any real attitude change. Second, attrition is

Table 21: Within-Kakuma hosting attitudes over time, repeated cross-section

	Oppose Camp Closure	Local Hosting	Nation Hosting
Wave 2 vs Wave 1	-0.332* (0.140)	-0.111 (0.123)	-0.275* (0.128)
Wave 3 vs Wave 2	0.057 (0.143)	0.042 (0.166)	0.629*** (0.155)
Age	-0.004 (0.010)	0.017† (0.009)	0.015 (0.009)
Gender: Male	-0.063 (0.171)	0.149 (0.160)	0.238 (0.160)
Education: primary	0.141 (0.285)	0.344 (0.263)	0.334 (0.278)
Education: secondary	0.497† (0.263)	0.441† (0.226)	0.667** (0.253)
Education: tertiary/adult	0.558* (0.273)	0.606* (0.243)	0.715** (0.257)
Constant	3.220*** (0.414)	2.178*** (0.372)	2.117*** (0.395)
Respondent-wave obs.	722	734	735
Respondents	301	302	302

Note:

OLS on the full repeated cross-section of Kakuma host-community respondents (at least 300 per wave, including refreshers), standard errors clustered by respondent in parentheses. Gender and education enter as indicators, with reference categories female and no formal education. Wave 3 vs Wave 2 is re-referenced to Wave 2; control coefficients and the constant come from the Wave 1-referenced models. Five-point outcomes, higher values more pro-refugee. *** $p < 0.001$, ** $p < 0.01$, * $p < 0.05$, † $p < 0.10$.

Table 22: Within-Kakuma hosting attitudes over four waves, balanced panel

	Oppose Camp Closure	Local Hosting	Nation Hosting
Wave 2 vs Wave 1	-0.481* (0.232)	0.169 (0.209)	0.024 (0.226)
Wave 3 vs Wave 2	0.073 (0.241)	-0.331 (0.243)	0.325 (0.212)
Wave 4 vs Wave 3	0.336 (0.270)	-0.030 (0.258)	0.525 (0.399)
Respondent-wave obs.	330	331	257
Respondents	83	83	83

Note:

OLS with respondent fixed effects on the balanced four-wave panel; standard errors clustered by respondent, in parentheses. Rows are consecutive wave-on-wave changes. Wave 4 is the February 2026 phone wave. Higher values indicate more pro-refugee attitudes. *** $p < 0.001$, ** $p < 0.01$, * $p < 0.05$, † $p < 0.10$.

substantially higher than across the three in-person waves: of the 120 Wave 4 respondents, 83 are present in all four waves. To hold sample composition fixed, Table 22 restricts to these 83 respondents, tracing within-respondent change in support for local hosting and opposition to camp closure from May 2025 through February 2026.

Across Waves 3 and 4 none of the changes are statistically significant. Descriptively, opposition to camp closure continues to recover, nation hosting rises further, and local hosting is essentially flat. The pattern is consistent with a partial recovery, or at least no further backlash, as aid was partly reinstated. However, with only 83 respondents these estimates are underpowered, and I read the trajectories as suggestive.

3.5 Integration and trust measures

The main text shows that, over the period of the cuts, support for closing the camp rose in Kakuma even as support for hosting refugees in Kenya more broadly did not fall. This points to a distinction between opposition to hosting refugees locally and opposition to their presence in the country more generally: a host can become less willing to have refugees in their own area without opposing their country providing asylum at all. To probe this distinction further, I examine four additional panel measures that capture attitudes toward refugees beyond local hosting: support for three integration rights (access to healthcare, the right to work, and freedom of movement) and a five-point measure of trust in refugees. Each is estimated with the same event-study specification and balanced three-wave sample as the main panel analysis.

Table 23 shows no decline in any of the four measures. None of the Wave 2 versus Wave 1 differences is distinguishable from zero, so neither support for integration nor trust in refugees fell during the trough of the aid contraction. If anything, both moved in the opposite direction over the stabilization period: support for freedom of movement rose between Waves 2 and 3 (+0.465, $p < 0.01$) and trust in refugees rose marginally (+0.229, $p < 0.10$), mirroring the recovery in nationwide hosting support documented in the main text. Consistent with the empathy channel and the open-ended evidence, the decline in willingness to host refugees locally was not accompanied by a broader decline in support for refugee rights or trust in refugees.

Table 23: Refugee integration and trust over the Kakuma panel

	Access to healthcare	Right to work	Freedom of movement	Trust in refugees
Wave 2 vs Wave 1	-0.198 (0.151)	-0.039 (0.116)	-0.093 (0.130)	-0.173 (0.121)
Wave 3 vs Wave 2	-0.053 (0.155)	0.132 (0.145)	0.465** (0.144)	0.229† (0.117)
Respondent-wave obs.	609	615	614	606
Respondents	205	205	205	205

Note: OLS with respondent fixed effects on the balanced three-wave Kakuma panel; standard errors clustered by respondent, in parentheses. Wave 3 vs Wave 2 is re-referenced to Wave 2. Five-point outcomes, higher values more supportive. *** $p < 0.001$, ** $p < 0.01$, * $p < 0.05$, † $p < 0.10$.

3.6 Personal-loss difference-in-differences

3.6.1 Comparing baseline-beneficiaries with non-baseline-beneficiaries

Because selection into aid benefits is not random, Table 24 compares baseline-beneficiaries and non-baseline-beneficiaries. Baseline-beneficiaries ($N = 127$) are demographically similar to non-baseline-beneficiaries ($N = 78$) on gender and education, and marginally older (+2.45 years, $p = 0.05$). They differ, however, on pre-shock attitudes: at Wave 1, baseline-beneficiaries were substantially more supportive of hosting nationally (+0.73, $p < 0.01$) and locally (+0.93, $p < 0.001$), and substantially more opposed to closing the camp (−0.95, $p < 0.001$).

This difference across the two populations was expected and in-line with the paper’s theory. Households tied to the refugee-aid economy plausibly hold more pro-refugee views both because they benefit from that economy and because pre-existing support may have increased their ability or desire to work and access services alongside refugees. The identifying assumption of the ITT specification is not baseline balance but

parallel trends: the respondent fixed effects absorb these level differences, and the estimates in Table 25 identify the differential change in exposed households’ attitudes after the cuts, not the difference in their levels. The remaining threat is differential trends, that exposed households would have converged toward unexposed households’ attitudes even absent the cuts. The Wave 2 placebo in the cohort specification (Table 25, Panel B) speaks against this, and the triangulation across the three H2 designs in Section 3.6.2 addresses it further.

Table 24: ITT balance: baseline-beneficiary vs non-baseline-beneficiary

Variable	Baseline exposed (Mean, SD)	Unexposed (Mean, SD)	Difference	p-value
Age (years)	31.10 (9.72)	28.65 (7.92)	2.45†	0.051
Female (%)	0.47 (0.50)	0.51 (0.50)	-0.04	0.577
Education index	0.65 (0.48)	0.71 (0.46)	-0.05	0.443
Hosting support — national (1–5)	3.40 (1.71)	2.67 (1.78)	0.73**	0.004
Hosting support — local (1–5)	3.46 (1.70)	2.54 (1.73)	0.93***	0.000
Support for camp closure (1–5)	2.25 (1.64)	3.21 (1.88)	-0.95***	0.000
Empathy index (0–1, measured W3)	0.66 (0.35)	0.62 (0.39)	0.04	0.406

Note: Wave 1 characteristics by whether the household was an aid recipient or NGO/aid employee at baseline. Columns report mean and standard deviation for the treated (exposed) and control (unexposed) groups; Difference is treated minus control. Significance: *** $p < .001$, ** $p < .01$, * $p < .05$, † $p < .10$ (Welch’s t -test).

3.6.2 Alternative personal loss measures

The main text reports the baseline-beneficiary (intent-to-treat) specification as the primary test of H2, chosen because its treatment is fixed before the cuts and therefore cannot be contaminated by later attitudes. Here I report all three of the personal-loss specifications I estimated, which trade identification against power and define exposure in three different ways. Selection into loss is non-random in every case, so I read all three as quasi-experimental evidence on the egocentric channel rather than as definitive estimates of the effect of personal loss. Their value lies in triangulation: each has a different weakness, and they point the same way. The three differ in how they define which households are exposed to personal loss, not in the estimator, which is a respondent-fixed-effects difference-in-differences throughout (equations (4) through (6)).

The first is the *baseline-beneficiary* specification (Panel A), an intent-to-treat design. Treatment is a fixed Wave 1 trait: whether the household received aid or contained an NGO or aid employee before the severe cuts. Because this indicator is measured before the cuts, it cannot be a product of how a household later came to feel about refugees, which makes it the most credible of the three on identification grounds. The cost is a weaker effect estimate because not every baseline beneficiary actually lost a benefit, so the “treated” group includes households that were not affected, which pulls the estimated effect downward. That the estimate remains significant despite this strengthens rather than weakens the inference.

One parameterization note applies to Panel A. The main text’s Table 3 reports the intent-to-treat specification in pooled form (equation (5)), replacing the respondent fixed effects with a baseline-beneficiary main effect so that the pre-cut level gap between beneficiary and non-beneficiary households is displayed alongside the interactions. Panel A here re-estimates the same design with respondent fixed effects (equation (4)), matching the estimator used in Panels B and C. In a fully balanced panel the two parameterizations return identical interaction estimates; they differ here only through item-level nonresponse on the outcome variables, and the camp-closure interactions are nearly identical across the two.

The second is the *cohort difference-in-differences* (Panel B), which defines exposure by realized loss. It compares households that first lost access to a benefit between Waves 2 and 3 against households that never lost access across the three waves, using the Wave 1 to Wave 2 interval as a placebo. Because loss is realized rather than assigned, this is closer to a treatment-on-treated estimate and is more exposed to the concern that attitudes drive loss classification. Its strength is the explicit falsification test: the Wave 2 interaction is a placebo that should be near zero if the two groups were on parallel trends before the Wave 3 shock, isolating the effect of the loss itself.

The third is the *switching difference-in-differences* (Panel C), which defines exposure by within-household loss events. It pools all loss events across waves and domains, regressing the outcome on an indicator that switches on once a household has lost access in a domain, with respondent and wave fixed effects. This design has the most power and, because it estimates a separate coefficient for each domain, shows whether the result is driven by one channel or holds across education, healthcare, refugee employment, and NGO employment. Its weakness is that loss is measured post treatment and that it has no accompanying placebo test.

Table 25: Within-Kakuma hosting attitudes by household exposure to the cuts

	Oppose Camp Closure	Local Hosting	Nation Hosting
<i>Panel A: Intent-to-treat, baseline beneficiary status</i>			
Baseline beneficiary $\times$ Wave 2	-0.711* (0.311)	0.067 (0.269)	-0.102 (0.289)
Baseline beneficiary $\times$ Wave 3	-0.712* (0.336)	-0.676† (0.352)	-0.603† (0.335)
Respondent-wave obs.	612	613	614
Respondents	204	205	205
<i>Panel B: Cohort DiD, newly lost W2–W3 vs never lost</i>			
Newly lost W2→W3 $\times$ Wave 2 (placebo)	-0.493 (0.389)	-0.165 (0.342)	0.171 (0.357)
Newly lost W2→W3 $\times$ Wave 3	-0.840* (0.398)	-0.577 (0.466)	-0.562 (0.425)
Respondent-wave obs.	375	376	377
Respondents	125	126	126
<i>Panel C: Switching DiD by loss domain</i>			
Any domain	-0.606* (0.240)	-0.325 (0.238)	-0.583** (0.222)
Education	-0.311 (0.342)	-0.030 (0.290)	-0.207 (0.307)
Healthcare	-0.790** (0.272)	-0.024 (0.257)	-0.357 (0.224)
Refugee employment	-0.970* (0.440)	-0.372 (0.383)	-0.655 (0.399)
NGO employment	-0.629* (0.262)	-0.494† (0.293)	-0.321 (0.292)
Respondent-wave obs.	612	613	614
Respondents	204	205	205

Note: OLS difference-in-differences on the balanced three-wave Kakuma panel; standard errors clustered by respondent in parentheses. Panels A and B interact an exposure indicator with wave indicators (Wave 1 omitted), with respondent fixed effects; Panel C regresses each outcome on an indicator switching on from the wave after the household first loses access, with respondent and wave fixed effects, each domain a separate model. Exposure is baseline-beneficiary status in Panel A (camp healthcare or education, or NGO or refugee employment, in the 30 days before Wave 1; 127 beneficiary, 78 non-beneficiary households); first loss between Waves 2 and 3 versus never losing access in Panel B, excluding Wave 1–2 losers; and realized loss events pooled across waves in Panel C. Five-point outcomes, higher values more pro-refugee. Significance: *** $p < .001$, ** $p < .01$, * $p < .05$, † $p < .10$.

Table 25 reports all three specifications across the two local outcomes and nationwide hosting. The three converge on the same conclusion for *Oppose Camp Closure*: households exposed to personal loss moved toward more pro-closure attitudes relative to their unexposed neighbors. The baseline-beneficiary specification shows exposed households declining at both the trough and the partial-recovery wave ($\hat{\tau}_2 = -0.711$, $\hat{\tau}_3 = -0.712$, both $p < 0.05$). The cohort difference-in-differences finds a larger Wave 3 effect among new-loss households ($\hat{\tau}_3 = -0.840$, $p < 0.05$), with a Wave 2 placebo that is insignificant, consistent with parallel pre-trends. The switching specification gives an any-domain estimate of -0.606 ($p < 0.05$), with consistent signs across all four domains and significant declines for healthcare (-0.790 , $p < 0.01$), refugee employment (-0.970 ,

$p < 0.05$), and NGO employment (-0.629 , $p < 0.05$); only education is individually null. The pattern of declining support among those who personally lose out holds across all three despite their different identifying assumptions.

Table 26: Intent-to-treat estimates with inverse-probability-of-retention weights

	Oppose Camp Closure	Local Hosting	Nation Hosting
<i>Panel A: Unweighted (main specification)</i>			
Baseline beneficiary $\times$ Wave 2	-0.711* (0.311)	0.067 (0.269)	-0.102 (0.289)
Baseline beneficiary $\times$ Wave 3	-0.712* (0.336)	-0.676† (0.352)	-0.603† (0.335)
Respondent-wave obs.	612	613	614
<i>Panel B: Inverse-probability-of-retention weights</i>			
Baseline beneficiary $\times$ Wave 2	-0.670* (0.308)	0.094 (0.268)	-0.062 (0.290)
Baseline beneficiary $\times$ Wave 3	-0.666† (0.338)	-0.608† (0.352)	-0.560† (0.335)
Respondent-wave obs.	612	613	614
<i>Note:</i> Respondent-fixed-effects intent-to-treat specification (baseline beneficiary status interacted with wave indicators) on the balanced three-wave Kakuma panel, standard errors clustered by respondent in parentheses. Panel B weights each retained respondent by the normalized inverse of the fitted probability of appearing in all three waves, from a logit of retention on Wave 1 age, gender, education, hosting support, and beneficiary status. Significance: *** $p < .001$, ** $p < .01$, * $p < .05$, † $p < .10$.			

4 Survey experiment results and robustness

4.1 Treatment arm balance

Each respondent rated six vignette profiles whose four attributes — the level of aid, whether services are shared with hosts, whether organizations hire locally, and whether conflict is ongoing — were independently randomized. If randomization succeeded, attribute assignment should be unrelated to respondents’ pre-treatment characteristics. Table 27 reports, for each attribute, the share of profiles at each level, and tests whether that assignment is predicted by respondent age, gender, or education. Each column is a regression of an attribute-level indicator on the three covariates; the reported statistic is the p -value from a joint test that all three coefficients are zero, with standard errors clustered by respondent. The results show that no attribute is jointly predicted by the covariates, consistent with successful randomization.

4.2 Full AMCE estimation results

Table 28 reports the full estimation results underlying Figure 4 of the main paper. The model is a linear regression of rated hosting support on indicators for each non-reference attribute level, with respondent fixed effects and standard errors clustered by respondent. Every attribute-level coefficient shown in the figure appears here alongside its standard error, sample size, and dependent variable. The estimates differ slightly from the “at mean empathy” column of Table 31, which comes from a separate model that interacts each attribute with the empathy index.

4.3 Comprehension check

The vignette instrument included a comprehension check, pre-specified in the pre-analysis plan, asking respondents the level of food and cash support in the last scenario they rated. Comprehension was high: 97.7 percent of respondents answered correctly, and the pass rate was stable across the three aid levels (97.8

Table 27: Vignette treatment-arm balance

Level	Share	Profiles	Joint p
Aid level (ref: 100%)			
40%	0.32	1,809	0.57
0%	0.33	1,809	0.75
Service access (ref: shared)			
Refugees only	0.51	1,809	0.14
Local hiring (ref: hire locals)			
Do not hire locals	0.50	1,809	0.55
Conflict status (ref: ongoing)			
Peace restored	0.50	1,809	0.93

Note: One row per non-reference attribute level. Share is the proportion of rated profiles at that level; Profiles is the number of profile ratings. Joint p is a Wald test that respondent age, gender, and education do not jointly predict assignment to that level, from a linear probability model with standard errors clustered by respondent. Respondents with missing demographics are excluded.

Table 28: Vignette experiment: full AMCE estimation results

	Nation Hosting
Aid: 40% (vs 100%)	-0.138† (0.083)
Aid: 0% (vs 100%)	-0.189* (0.080)
Hiring: do not hire locals (vs hire)	-0.203** (0.072)
Access: refugees only (vs shared)	-0.184* (0.076)
Conflict: peace restored (vs ongoing)	-0.633*** (0.081)
Respondent-wave obs. (profiles)	1,883
Respondents	299

Note: Average marginal component effects from a linear model with respondent fixed effects; standard errors clustered by respondent in parentheses. The outcome is rated support for Kenya hosting refugees (1–5) under each vignette (the *Nation Hosting* measure). Reference levels are aid at 100 percent, hire local residents, services shared with hosts, and conflict ongoing. Estimated separately from the empathy-interaction model in Table 31. *** $p < 0.001$, ** $p < 0.01$, * $p < 0.05$, † $p < 0.10$.

percent at 100 percent aid, 97.0 at 40 percent, and 98.7 at 0 percent). Table 29 re-estimates the average marginal component effects excluding those who failed, alongside the full-sample estimates. The estimates are essentially unchanged. Because comprehension was nearly universal and the results do not depend on excluding the few who failed the check, I retain the full sample in the main analysis.

4.4 Hypothetical full-cut scenario

In Waves 1 and 2, I asked panel respondents how strongly they would support Kenya continuing to host refugees if international donors withdrew financial support completely, a more severe cut than the one

Table 29: Vignette AMCEs excluding comprehension-check failures

	Full sample	Passed comprehension check
Aid: 40% (vs 100%)	-0.138† (0.083)	-0.148† (0.085)
Aid: 0% (vs 100%)	-0.189* (0.080)	-0.188* (0.082)
Hiring: do not hire locals (vs hire)	-0.203** (0.072)	-0.188** (0.073)
Access: refugees only (vs shared)	-0.184* (0.076)	-0.197* (0.077)
Conflict: peace restored (vs ongoing)	-0.633*** (0.081)	-0.604*** (0.080)
Respondent-wave obs. (profiles)	1,883	1,853
Respondents	299	294

Note: Average marginal component effects from linear models with respondent fixed effects; standard errors clustered by respondent in parentheses. The outcome is rated support for Kenya hosting refugees (1–5) under each vignette. The second column excludes respondents who failed the comprehension check. *** $p < 0.001$, ** $p < 0.01$, * $p < 0.05$, † $p < 0.10$.

Table 30: Hypothetical full-cut scenario: within-respondent change and gap to general hosting support

	Support
<i>Panel A: Over-time change in full-cut support</i>	
Wave 2 vs Wave 1 (5-point)	-0.545*** (0.150)
Wave 2 vs Wave 1 (support ≥ 4)	-0.136*** (0.040)
Respondent-wave obs.	396
<i>Panel B: Gap to general hosting support</i>	
Full-cut vs. general (Wave 1 gap)	-0.473** (0.153)
Full-cut $\times$ Wave 2	-0.367† (0.199)
Respondent-wave obs.	813

Note:

OLS with respondent fixed effects on the balanced Wave 1–Wave 2 panel; standard errors clustered by respondent, in parentheses. The outcome is support for hosting refugees if donors withdrew aid completely, five-point (higher = more supportive); the binary row codes support as four or five. Panel A is the within-respondent Wave 1 to Wave 2 change. Panel B stacks the full-cut item against general hosting support, giving the Wave 1 gap and its change by Wave 2. *** $p < 0.001$, ** $p < 0.01$, * $p < 0.05$, † $p < 0.10$.

observed. Because the item is within-respondent and repeated across the two waves that bracket the trough of the contraction, I use it in three ways: the level of support gauges attitudes under a hypothetical full withdrawal at each wave; the gap between this item and general support for hosting refugees in Kenya shows how attitudes differ within respondent under current conditions versus a complete cut; and the within-respondent change from Wave 1 to Wave 2 asks whether support fell as the cuts deepened. This item is

a descriptive counterpart to the full-cut arm of the vignette experiment. Standard errors are clustered by respondent throughout, and these estimates are descriptive.

The findings support the main experimental result. At baseline, support under full withdrawal was 0.473 points lower than support for hosting refugees in Kenya generally ($p < 0.01$), and this gap widened by a further 0.367 points by Wave 2 ($p < 0.10$) (Table 30). Support for hosting under full withdrawal also fell sharply across waves, from 2.64 to 2.12 on the five-point scale ($\hat{\beta} = -0.545$, $p < 0.001$), while the share expressing support fell from 42.1 to 29.4 percent ($p < 0.001$). As the real cuts deepened and the prospect of hosting without aid became more concrete, opposition to this scenario grew. Yet even under hypothetical complete withdrawal, 29 percent still supported hosting, indicating that attitudes do not rest on aid alone.

4.5 Humanitarian concern and empathy

The qualitative and experimental findings point to a possible role for empathy that was not pre-specified. Respondents frequently framed their views in humanitarian terms in the open-ended questions, and the *Conflict* attribute in the vignette had a larger effect on hosting support than any of the aid-related attributes. I therefore examine whether more empathetic respondents react differently to reductions in aid.

I measure empathy using a standardized index that averages two five-point items: concern about refugees' suffering and worry that refugees lack sufficient food because of the cuts. I interact this index with each attribute in the vignette regression, including respondent fixed effects and clustering standard errors by respondent. Because the fixed effects absorb the main effect of empathy, the interaction coefficients show how each attribute contrast changes with a one-standard-deviation increase in empathy.

Table 31 reports the results. The interaction between empathy and the 40 percent aid contrast is positive and significant (0.230, $p < 0.01$). At mean empathy, reducing aid to 40 percent lowers hosting support by 0.137 points, but a one-standard-deviation increase in empathy offsets nearly all of this decline. The interaction for a complete withdrawal of aid is also positive, though smaller and less precisely estimated (0.131, $p \approx 0.10$). Empathy does not significantly condition responses to local hiring. For shared services, the interaction is negative and marginally significant; Section 4.5.2 shows that this reflects the higher level of support among high-empathy respondents when services are shared with hosts.

The descriptive results show the same pattern. Among respondents below the median of the empathy index, support for hosting falls from 40.9 percent under full aid to 32.5 percent under 40 percent aid and 33.2 percent under no aid. Among those above the median, support remains almost unchanged: 40.3, 40.3, and 39.8 percent, respectively. The two groups begin at nearly identical levels under full aid, but only the low-empathy group withdraws support as aid falls.

These results are exploratory and should not be interpreted causally. Empathy was not randomly assigned, was measured alongside the experiment rather than beforehand, and the high-empathy subgroup is relatively small. The analysis was also not preregistered, having emerged from the open-ended responses and refugee-leader interviews. I therefore treat the findings as suggestive.

4.5.1 Empathy index construction and robustness

The vignette heterogeneity analyses use a two-item empathy index averaging concern about refugees and their suffering with worry that refugees lack enough food because of the cuts. A third item, worded in the opposite direction (that the respondent rarely feels compassion), was fielded alongside these two and reverse-coded to the common orientation. I exclude it from the primary index because it does not cohere with the other two: Table 32 reports the pairwise correlations, which are strong and positive between the two retained items but negative between the recoded compassion item and each of them. The two retained items form an internally consistent pair (Cronbach's $\alpha = 0.72$), whereas adding the compassion item yields a negative alpha (-0.31), indicating that the three items do not scale together and that the compassion item cannot be treated as a third measure of the same construct.

Table 33 shows that the substantive conclusion does not depend on this choice: re-estimating equation (8) with the three-item index in place of the two-item index leaves the sign and pattern of the Aid interactions

Table 31: Vignette experiment by empathy

	Nation Hosting	
	AMCE at mean empathy	$\times$ Empathy (interaction)
Aid: 40% (vs 100%)	-0.137 [†] (0.082)	0.230** (0.080)
Aid: 0% (vs 100%)	-0.180* (0.080)	0.131 (0.080)
Hiring: Do not hire locals (vs hire)	-0.211** (0.073)	0.046 (0.074)
Access: Refugees only (vs shared)	-0.187* (0.076)	-0.136 [†] (0.077)
Conflict: Peace restored (vs ongoing)	-0.623*** (0.080)	0.030 (0.084)

Note: Linear model with respondent fixed effects. Standard errors (in parentheses) clustered by respondent. Empathy is the two-item index, standardized (mean 0, SD 1), so the AMCE column gives each contrast at mean empathy and the interaction column the change in that contrast per one-standard-deviation increase in empathy. $N = 1,885$ profiles from 301 respondents. *** $p < 0.001$, ** $p < 0.01$, * $p < 0.05$, [†] $p < 0.10$.

unchanged, with the Aid 40% interaction attenuating slightly (0.230 to 0.220) and the Aid 0% interaction essentially unchanged (0.131 to 0.141). Because neither the empathy channel nor the Aid $\times$ empathy interaction was pre-registered, these analyses are exploratory and I read them as suggestive rather than confirmatory.

Table 32: Empathy items: pairwise correlations and internal consistency

	Item 1	Item 2	Item 3
Concern for refugees' suffering (item 1)			
Worry about refugees' food (item 2)	0.56		
Reverse-coded item (item 3)	-0.40	-0.39	
Cronbach's alpha: two-item / three-item	0.72	-0.31	

Note: Pairwise Pearson correlations among the three empathy items on the vignette respondents with all three items observed ($N = 301$). All items are oriented so that higher values indicate more empathy; item 3 is reverse-coded and recoded. The alpha row reports Cronbach's alpha for the two-item index (items 1–2) and the three-item index.

Table 33: Empathy moderation under the two-item and three-item indexes

	$\times$ Empathy (two-item index)	$\times$ Empathy (three-item index)
Aid: 40% (vs 100%)	0.230** (0.080)	0.220** (0.079)
Aid: 0% (vs 100%)	0.131 (0.080)	0.141 [†] (0.078)
Hiring: Do not hire locals (vs hire)	0.046 (0.074)	0.058 (0.084)
Access: Refugees only (vs shared)	-0.136 [†] (0.077)	-0.088 (0.073)
Conflict: Peace restored (vs ongoing)	0.030 (0.084)	0.084 (0.093)

Note: Attribute-by-empathy interaction coefficients from linear models with respondent fixed effects, standard errors clustered by respondent in parentheses. Each column standardizes the stated empathy index to mean zero and standard deviation one across respondents. The three-item column adds the reverse-coded item to the index. N (two-item) = 1,883 profiles; N (three-item) = 1,883 profiles. Significance: *** $p < 0.001$, ** $p < 0.01$, * $p < 0.05$, [†] $p < 0.10$.

4.5.2 Empathy marginal means

Interaction terms in equation (8) compare conditional AMCEs, which depend on the reference level of each attribute. Following Leeper, Hobolt, and Tilley (2020), subgroup comparisons in factorial designs are more

transparently reported as marginal means, the average rated support at each attribute level within each subgroup. Table 34 reports marginal means for every attribute level separately for respondents below and above the sample median of the empathy index, with standard errors clustered by respondent, and the high-minus-low difference at each level.

The marginal means reproduce the interaction results. Among low-empathy respondents, support falls as aid drops from 100 percent to 40 and 0 percent, while among high-empathy respondents support is essentially flat across aid levels; the high-minus-low differences are close to zero at full aid and positive under the cut conditions.

Table 34: Vignette experiment: marginal means by empathy group

Level	Low empathy	High empathy	Difference (High – Low)
<i>Aid level</i>			
100%	2.640 (0.105)	2.629 (0.114)	-0.012 (0.155)
40%	2.356 (0.100)	2.572 (0.109)	0.216 (0.148)
0%	2.386 (0.103)	2.599 (0.113)	0.212 (0.153)
<i>Local hiring</i>			
Hire locals	2.581 (0.094)	2.686 (0.108)	0.104 (0.143)
Do not hire locals	2.349 (0.088)	2.512 (0.095)	0.162 (0.129)
<i>Service access</i>			
Shared with hosts	2.480 (0.095)	2.715 (0.099)	0.235† (0.137)
Refugees only	2.448 (0.095)	2.485 (0.100)	0.037 (0.138)
<i>Conflict status</i>			
Conflict ongoing	2.751 (0.100)	2.892 (0.101)	0.141 (0.142)
Peace restored	2.187 (0.095)	2.298 (0.098)	0.111 (0.136)

Note:

Marginal means of rated hosting support (1–5) at each attribute level, by empathy group (median split on the two-item empathy index), with standard errors clustered by respondent in parentheses. The difference column reports the high-minus-low contrast among profiles at that level. Low-empathy group: 158 respondents; high-empathy group: 143 respondents. Significance (difference column): *** $p < .001$, ** $p < .01$, * $p < .05$, † $p < .10$.

5 Additional analyses and study documentation

5.1 Enumerator fixed effects

TIFA provided interviewer identifiers for the three Kakuma panel waves, the February 2026 nationwide survey, and the May 2026 recontact. Identifiers do not exist for the 2023 baseline. Table 35 re-estimates the personal-loss specifications of Table 25 with and without enumerator fixed effects. Table 35 shows that including enumerator fixed effects leaves the estimates substantively unchanged, suggesting that the declines in hosting support among personally exposed households is not an artifact of who conducted the interview.

I do not report an analogous check for the over-time wave coefficients: 61 percent of Wave 3 interviews were conducted by enumerators who did not work the first two waves, so enumerator fixed effects are nearly collinear with the Wave 3 indicator. For those estimates, interviewer composition is instead addressed by the repeated cross-section (Section 3.3), which recovers the over-time pattern on refreshed samples with a different enumerator mix at each wave.

The 2023 baseline data does not include enumerator identifiers, so I cannot examine enumerator effects across the full 2023–2026 difference-in-differences. But because the interaction is the 2026 host-county gap net of the fixed 2023 gap, showing that the 2026 gap is stable to enumerator changes provides indirect evidence that the design is not driven by interviewer effects. Table 36 re-estimates the 2026 host-county comparison with and without enumerator fixed effects, using the same controls and standard errors as the difference-in-difference. The fixed effects are identified because 31 of 36 February enumerators interviewed

Table 35: Within-Kakuma personal-loss estimates with enumerator fixed effects

	Oppose Camp Closure		Local Hosting		Nation Hosting	
	Base	Enum. FE	Base	Enum. FE	Base	Enum. FE
<i>Panel A: Intent-to-treat, baseline beneficiary status</i>						
Beneficiary $\times$ Wave 2	-0.711* (0.311)	-0.719* (0.317)	0.067 (0.269)	0.070 (0.271)	-0.102 (0.289)	-0.177 (0.288)
Beneficiary $\times$ Wave 3	-0.712* (0.336)	-0.826** (0.318)	-0.676† (0.352)	-0.695* (0.351)	-0.603† (0.335)	-0.712* (0.328)
Respondent-wave obs.	612	612	613	613	614	614
Enumerators		18		18		18
<i>Panel B: Cohort DiD, newly lost W2–W3 vs never lost</i>						
Newly lost $\times$ Wave 2 (placebo)	-0.493 (0.389)	-0.501 (0.397)	-0.165 (0.342)	-0.203 (0.358)	0.171 (0.357)	0.241 (0.361)
Newly lost $\times$ Wave 3	-0.840* (0.398)	-0.669† (0.405)	-0.577 (0.466)	-0.509 (0.509)	-0.562 (0.425)	-0.551 (0.452)
Respondent-wave obs.	375	375	376	376	377	377
Enumerators		18		18		18
<i>Panel C: Switching DiD, any domain</i>						
Lost aid (post)	-0.606* (0.240)	-0.554* (0.242)	-0.325 (0.238)	-0.242 (0.242)	-0.583** (0.222)	-0.603** (0.229)
Respondent-wave obs.	612	612	613	613	614	614
Enumerators		18		18		18

Note: Same estimators, samples, and outcomes as Table 25, restricted to interviews with a matched enumerator identifier so paired columns share a sample. Base columns absorb respondent fixed effects (and wave fixed effects in Panel C); Enum. FE columns additionally absorb the wave-specific interviewer. Over-time wave coefficients are omitted, being collinear with enumerator identity. Standard errors clustered by respondent in parentheses. Significance: *** $p < .001$, ** $p < .01$, * $p < .05$, † $p < .10$.

Table 36: Nationwide 2026 host-county gap with enumerator fixed effects

	Oppose Repat. (Feb)		Local Hosting (Feb)		Nation Hosting (May)	
	Base	Enum. FE	Base	Enum. FE	Base	Enum. FE
Host county	0.009 (0.023)	0.008 (0.024)	-0.290** (0.100)	-0.308** (0.103)	0.137 (0.171)	0.107 (0.173)
Observations	2,293	2,293	2,306	2,306	1,769	1,769
Enumerators		34		34		21

Note:

OLS on the 2026 sample, unweighted, with age, gender and education controls and HC1 standard errors in parentheses. Host county indicates Turkana or Garissa. Each pair of columns is restricted to respondents with a non-missing outcome and enumerator identifier. The first two outcomes come from the February survey and the nationwide hosting measure from the May recontact. Significance: *** $p < .001$, ** $p < .01$, * $p < .05$, † $p < .10$.

respondents in both host and non-host counties. The *Local Hosting* gap, the post-period component that drives the difference-in-differences result, moves from -0.290 to -0.308 , and the repatriation and nationwide hosting gaps are indistinguishable from zero in both columns. Together with the panel results, this indicates that neither the within-Kakuma erosion nor the nationwide gap reflects which enumerators conducted the interviews.

5.2 Minimum detectable effects (MDE)

The Kakuma panel is small, so before reading any near-zero estimate as evidence of no effect, I report the minimum detectable effect (MDE) for each headline test. For a two-sided test at the 5% level with 80% power, $MDE = (z_{.975} + z_{.80}) \times SE = 2.80 \times SE$, computed from each model's realized clustered standard error. The MDE concern is most relevant for one near-zero estimate. The aggregate over-time change in

the hosting index (-0.015 SD, MDE 0.236 SD) rules out aggregate shifts larger than roughly a quarter of a standard deviation, so it is best read as the absence of a large aggregate movement rather than of any movement. The other headline effects are detected rather than null, so the MDE is not a concern.

Table 37: Minimum detectable effects for the headline tests

Test	N	Est.	Clustered SE	MDE (80%)
Over-time panel: W3 vs W1 (hosting index)	615	-0.015	0.084	0.236
ITT: baseline beneficiary $\times$ W3 (camp closure)	612	-0.712	0.336	0.942
Vignette AMCE: Aid 0% vs 100% (1–5)	1,883	-0.189	0.080	0.226

Note: Each row reports the minimum detectable effect (MDE) for one headline test, computed as 2.80 times the model’s clustered standard error (two-sided, 5% level, 80% power). Est. and Clustered SE are the point estimate and standard error from the corresponding results table. Units are native to each specification: standard-deviation units for the hosting index, scale points for camp-closure support (1–5), and support points for the vignette outcome (1–5).

5.3 Shirika Plan

The nationwide difference-in-differences spans a three-year window, over which attitudes could have shifted for reasons unrelated to the cuts. The most plausible such confounder is the Shirika Plan, Kenya’s 2021 framework (with a 2024 implementation plan) to transition refugee camps into integrated settlements. If hosts became aware of a policy to integrate refugees locally, that awareness could move hosting support independently of the aid cuts. Tables 38 and 39 assess this by comparing hosting support between respondents aware and unaware of the Plan, within survey round, nationwide and in the Kakuma panel.

Two features of the awareness data make the Shirika Plan an unlikely driver of the results. First, nationwide awareness is too low to move the aggregate: it rose over the study period but remained limited in both rounds (1.5 percent in 2023 and 11.4 percent in 2026), so the large majority of respondents had not heard of the Plan in either year. Among those who had, national hosting support was higher than among the unaware in both rounds, and local hosting support in 2026 was slightly lower, though not significantly so. Second, in the Kakuma host area awareness was already widespread and stable across waves (66.7 percent at Wave 3), so it cannot track the within-area decline documented in the main text. At Wave 3, respondents aware of the Plan were if anything more supportive of hosting and less supportive of closing the camp than the unaware. The Shirika Plan is therefore unlikely to account for the results: nationally it is too rare to shift the aggregate, and locally it is neither new nor associated with lower support.

Table 38: Nationwide support for hosting by awareness of the Shirika Plan

Outcome	Year	Aware	Unaware	Diff.	<i>p</i>
National hosting	2023	3.49	3.06	+0.44	0.028
National hosting	2026	3.83	3.59	+0.24	0.060
Local hosting	2023	3.39	2.58	+0.80	0.010
Local hosting	2026	2.91	3.14	-0.23	0.110

Note: Comparisons are within survey year. Outcomes are five-point items (higher = more support). Means survey-weighted; difference tested with an unweighted Welch t-test. Awareness was 1.5% in 2023 and 11.6% in 2026.

5.4 Reversion to the mean

Both difference-in-differences designs compare groups that were initially more supportive of refugees. In the nationwide sample, refugee-hosting counties began about one-third of a standard deviation above the rest of the country on the *Hosting Index*. In the panel, baseline-beneficiary households were also more supportive

Table 39: Support for hosting by awareness of the Shirika Plan, Kakuma (Wave 3)

Outcome	Aware	Unaware	Diff.	<i>p</i>
National hosting	3.59	3.04	+0.54	0.037
Local hosting	3.14	2.82	+0.33	0.225
Close camp	2.72	3.25	-0.53	0.054

Note: Outcomes are five-point items; higher values indicate more support, except close camp, where higher indicates more support for closing the camp. Unweighted means; difference tested with an unweighted Welch t-test. Awareness in the Kakuma area was 66.7 percent at Wave 3.

of local hosting and less supportive of camp closure. If these initial differences reflected temporary shocks or measurement error, later convergence could be mistaken for an effect of the aid cuts. Three features of the results make this unlikely.

First, the baseline differences are consistent with prior evidence rather than unusually high observations. Research from Kenya and other refugee-hosting settings finds that communities and households connected to the aid economy tend to hold more favorable attitudes toward refugees, partly because hosting brings jobs, services, business opportunities, and infrastructure (Alix-Garcia et al., 2018; Zhou, Grossman and Ge, 2023; Kadigo and Maystadt, 2023; Baseler et al., 2025; Zhou et al., 2025; Demir, Mayda and Maystadt, 2025). MacDonald and Lichtenheld (2025) similarly find higher support for refugee integration in Kenya’s hosting counties in 2023. The initial gaps therefore match established patterns.

Second, reversion cannot explain the selectivity of the nationwide decline. Because the surveys use repeated cross-sections of different respondents, individual-level measurement error cannot drive the result; the concern is instead that county-level sampling error inflated the 2023 hosting-county mean. But such reversion would pull down every correlated hosting outcome measured in the same counties, in proportion to its initial gap. Instead, the decline is concentrated in local hosting, which falls by roughly two-thirds of a standard deviation, while support for hosting refugees in Kenya more broadly moves far less and support for the various integration and refugee-numbers measures does not decline at all. Because hosting counties began above the national average on all of these measures, regression to the mean should have produced declines across them. The fact that only local hosting changed suggests that the result reflects attitudes toward hosting refugees nearby rather than sampling error.

Third, the panel results follow the timing of actual losses. Beneficiary status is defined by economic ties to the aid system, not by baseline attitudes. Households that lost benefits between Waves 2 and 3 moved in parallel with other households before the loss and declined only afterward. The largest changes also occurred in the specific domains in which access was lost, while non-beneficiary households remained stable across waves. A general return toward the mean would not predict this timing or domain-specific pattern.

5.5 Deviations from the pre-analysis plan

This study was pre-registered before the first panel wave and before the trajectory of the aid cuts was known, with a February 2026 amendment adding the fourth phone wave and the 2026 nationwide follow-up. This section documents where the analysis departs from the pre-analysis plan (PAP) and why.

5.5.0.1 Dropping the two-town difference-in-differences. The PAP’s primary design compared Kakuma with Isiolo town, selected as a control for its similar economy and political trends, but lack of a refugee presence. During the study period, Isiolo experienced security and political issues unrelated to the aid cuts that made it unsuitable as a control case. The major crisis happened in June 2025, when a motion to impeach the Isiolo governor led to community-wide violence and insecurity, including stone-throwing,

teargas, and gunfire.¹ Field reports from the research team corroborated heightened insecurity during the study period.

These events likely affected the study’s outcomes, particularly perceived safety, crime victimization, and attitudes toward government. Any Kakuma–Isiolo comparison estimated over this period may therefore confound the effect of the aid cuts with the effects of the political crisis and violence. All three survey waves were completed, but I dropped the two-town design after fieldwork concluded, on the basis of the events documented above, because the identifying assumption had failed rather than because of any feature of the estimates. For transparency, this appendix reports both the contamination evidence and the pre-registered specification.

Figure 1 plots the two towns’ trajectories on the security outcomes and hosting support. The share of Isiolo respondents feeling unsafe at night rose 26 percentage points over the study period, from 33 to 59 percent, despite the absence of any aid shock in Isiolo, a steeper rise than in Kakuma itself. As a consequence, the pre-registered design returns a significant positive coefficient on relative safety in Kakuma (Table 40), which taken at face value would imply that the aid cuts improved security in the hosting area. On the attitude outcomes, the pre-registered contrast on the camp-closure item is negative and significant at Wave 3 (-0.532 , $p < .05$), directionally consistent with the within-Kakuma results in the main text, but given the contamination I do not interpret any of these coefficients as effects of the aid cuts.




Figure 1: Two-town trajectories on security outcomes and hosting support

5.5.0.2 Promotion of the nationwide comparison to the primary H1 design. With the pre-registered control unusable, I use the two nationally representative surveys bracketing the cuts as the primary test of community-level exposure, comparing refugee-hosting counties to the rest of Kenya. Both data sources were described in the PAP (January 2026 amendment). Their combination into a difference-in-differences design, the inclusion of Garissa alongside Turkana as treated counties, the entropy-balancing post-stratification weights, and the wild cluster bootstrap inference were not pre-specified and were adopted in response to the failure of the two-town design. The Kakuma panel over-time analysis is unchanged from the PAP apart from the loss of the control town.

5.5.0.3 Reorganization of hypotheses. The PAP specified five hypotheses across three outcome families: economic and security outcomes (PAP-H1), attitudes toward refugees (PAP-H2), and attitudes toward the national government, the United States, and the United Nations (PAP-H3 through H5). The paper instead tests two hypotheses, both concerning hosting attitudes. This narrowing reflects the paper’s more specific focus on attitudes toward refugee hosting. The economic and security outcomes show how the cuts affected host households. I now treat these results as descriptive because they establish the context for

¹“Isiolo Governor Abdi Guyo Impeached,” *Daily Nation*, June 27, 2025, <https://nation.africa/kenya/counties/isiolo/isiolo-governor-abdi-guyo-impeached--5096528>; “How Isiolo Governor Guyo Survived Impeachment on a Technicality,” *Daily Nation*, July 9, 2025, <https://nation.africa/kenya/news/politics/how-isiolo-governor-guyo-survived-impeachment-on-a-technicality-5111156>.

Table 40: Pre-registered two-town difference-in-differences (transparency only)

Outcome	Treated x W2	SE (W2)	Treated x W3	SE (W3)	N
Hosting support	-0.182	(0.198)	0.180	(0.220)	1,158
Effects index	0.058	(0.063)	-0.117†	(0.067)	1,169
Integration index	-0.088	(0.080)	0.143	(0.092)	1,169
Oppose camp closure	-0.249	(0.217)	-0.532*	(0.230)	1,161
Full rights	0.031	(0.044)	0.098†	(0.050)	1,170
Feels unsafe at night	-0.119*	(0.058)	-0.127*	(0.062)	1,166
Crime victim, past month	0.092†	(0.055)	0.016	(0.060)	1,170

Note: Balanced two-town panel (205 Kakuma and 185 Isiolo respondents observed in all three waves). Respondent fixed effects with wave indicators (Wave 1 omitted); coefficients are the differential change in Kakuma relative to Isiolo at each post wave. Standard errors clustered by respondent. The camp-closure outcome is reversed so higher values indicate opposition to closing the camps. The control town was contaminated during the study period, so these are not effects of the aid cuts. Significance: *** $p < .001$, ** $p < .01$, * $p < .05$, † $p < .10$.

the paper’s main question rather than test its central argument. Section 3.2 and this appendix report all preregistered economic and security outcomes. Attitudes toward the Kenyan government, the United States, and the United Nations address the different question of how citizens assign responsibility for the cuts. I examine those outcomes in a companion working paper on attribution and accountability. This current paper therefore focuses on PAP-H2, attitudes toward refugee hosting. I divide that hypothesis into H1 and H2 to distinguish community-level exposure from individual-level exposure.

5.5.0.4 The empathy analyses are exploratory. Neither the empathy channel nor the Aid $\times$ empathy vignette interaction was pre-registered, and neither is tested as a hypothesis. Both emerged inductively from the open-ended panel responses and refugee-leader interviews, in which respondents repeatedly framed their views in humanitarian terms. Because this pattern appeared prominently across the qualitative evidence, I examine whether it is also visible in the quantitative results, and I report those analyses as a discussion rather than a test. They should be read as exploratory rather than confirmatory. A related limitation, discussed in Section 4.5, is that the empathy moderator was measured post-treatment at Wave 3.

5.5.0.5 Within-treatment (H2) specifications. The baseline-beneficiary intent-to-treat design, which fixes exposure at Wave 1 aid-recipient or aid-employment status, follows the PAP’s pre-specified within-treatment design. The cohort difference-in-differences and the switching difference-in-differences reported in Section 3.6.2 were added after fieldwork as robustness checks and were not pre-specified.

5.5.0.6 Attrition and follow-up survey. The PAP specified a target of 300 respondents per town in each wave, allowing for expected attrition of 15–20%. Actual attrition was higher. The balanced three-wave Kakuma panel includes 205 respondents, while the refreshed cross-sections includes at least 300 respondents per wave. Section 1.3.2 discusses attrition in detail. The fourth phone wave was implemented as specified in the January 2026 amendment, but attrition was substantial after the switch in survey mode (results are reported in Section 3.4). In the 2026 nationwide survey, a technical issue affected the *Nation Hosting* question and required a May 2026 recontact with a nationally representative sample of 1,776 respondents. Sections 1.3.1 and 1.1 provide further details.

References

- Alix-Garcia, Jennifer, Sarah Walker, Anne Bartlett, Harun Onder, and Apurva Sanghi. 2018. “Do Refugee Camps Help or Hurt Hosts? The Case of Kakuma, Kenya.” *Journal of Development Economics* 130: 66–83.
- Anderson, Michael L. 2008. “Multiple Inference and Gender Differences in the Effects of Early Intervention: A Reevaluation of the Abecedarian, Perry Preschool, and Early Training Projects.” *Journal of the American Statistical Association* 103(484): 1481–1495.
- Baseler, Travis, Thomas Ginn, Robert Hakiza, Helidah Ogude-Chambert, and Olivia Woldemikael. 2025. “Can Redistribution Change Policy Views? Aid and Attitudes toward Refugees.” *Journal of Political Economy* 133(9): 2874–2917.
- Demir, Cansu, Anna Maria Mayda, and Jean-François Maystadt. 2025. “The Political Impact of Refugees in Africa.” Working paper.
- Hainmueller, Jens. 2012. “Entropy Balancing for Causal Effects: A Multivariate Reweighting Method to Produce Balanced Samples in Observational Studies.” *Political Analysis* 20(1): 25–46.
- Kadigo, Mark Marvin, and Jean-François Maystadt. 2023. “How to Cope with a Refugee Population? Evidence from Uganda.” *World Development* 169: 106293.
- Leeper, Thomas J., Sara B. Hobolt, and James Tilley. 2020. “Measuring Subgroup Preferences in Conjoint Experiments.” *Political Analysis* 28(2): 207–221.
- MacDonald, Mae, and Adam Lichtenheld. 2025. “Support for Refugee Integration in a Major Refugee-Hosting Country.” Working Paper No. 25-01, Immigration Policy Lab, Stanford University.
- Zhou, Yang-Yang, Guy Grossman, and Shuning Ge. 2023. “Inclusive Refugee-Hosting Can Improve Local Development and Prevent Public Backlash.” *World Development* 166: 106203.
- Zhou, Yang-Yang, Naijia Liu, Shuning Ge, and Guy Grossman. 2025. “Liberalizing Refugee Hosting Policies without Losing the Vote.” OSF Preprints, Version 4.